

Ideology as a Luxury Good in Occupational Choice

Turgut Keskintürk*

Duke University

July, 2026

[\[link to the latest version\]](#)

Abstract

Career aspirations are formed long before people enter the labor market, but whether and how political ideology is part of that formative process remains poorly understood. Using yearly surveys of 9.5 million college freshmen entering U.S. institutions between 1976 and 2010, I show that ideology systematically structures occupational aspirations. This pattern is two to three times larger among students from the highest socioeconomic backgrounds than among those from the lowest. To explain this difference, I propose a luxury-good mechanism: ideological congruence, the match between one's politics and an occupation's perceived ideological orientation, constitutes an occupational amenity whose importance increases with material security. This, I theorize, makes ideology an expressive resource for life choices, though one that is unequally distributed.

Food is the first thing. Morals follow on.

— Bertolt Brecht
The Threepenny Opera (1928)

Introduction

Workplace stratification by political ideology is a consequential feature of the contemporary American labor market ([Chinoy and Koenen 2024](#)), with some firms disproportionately liberal and others disproportionately conservative (see [Frake, Hurst, and Kagan 2023](#); [Kagan, Frake, and Hurst 2025](#)).

*I would like to thank Andrés Castro Araújo, Pablo Bello, Jennie Brand, Braulio Güémez, William Holtkamp, Michelle Jackson, Christopher Johnston, Kevin Kiley, Hedy Lee, Scott M. Lynch, Martín Opertti, Alejandro Sarria, Elizaveta Sheremet, Robin Stryker, Ang Yu, Kim Weeden, and Stephen Vaisey for their valuable comments on this project, and conference participants at Duke University and Stanford University for their helpful feedback. I sincerely thank Ellen Bara Stolzenberg and the Higher Education Research Institute for their help in acquiring this data, and I am grateful to the Institute for Humane Studies for their support for this project under grant no. IHS020598.
Correspondence: turgut.keskinturk@duke.edu.

Such political sorting is not confined to firms: it characterizes entire occupational groups, with partisan ratios as large as 4 to 1 (meaning four Democrats for every Republican and vice versa), within certain occupational categories.¹ Put differently, there is a fundamental *alignment* between the kind of work people do and the kind of political views they favor (Weeden and Grusky 2005).

Two central mechanisms, operating at different points in the life course, may explain such political sorting. On the one hand, people may form politically congruent occupational preferences through early-life socialization and self-select into ideologically consistent occupations (e.g., Meriläinen and Mitrunen 2025; also see Bermiss and McDonald 2018). This would establish political socialization as a foundational process that shapes major career decisions. On the other hand, ideological sorting may emerge after occupational entry, as people develop political views via workplace socialization (e.g., Kitschelt and Rehm 2014; Selenko et al. 2025; Weeden and Grusky 2005). This would position political ideology as a consequence of occupations, rather than a driver of prior choice.²

One significant challenge in distinguishing these two mechanisms is that existing research observes workers only after they have sorted into occupations. While this approach is valuable for precisely documenting the extent and magnitude of political sorting (Chinoy and Koenen 2024), the realized positions reflect the outcome of multiple processes: workplace-level sorting and preference-based self-selection, certainly, but also structural forces such as credentialing, residential segregation, and institutional gatekeeping (e.g., Becker [1964] 1994; Brown and Enos 2021; Weeden 2002). Since the political composition of occupations is the equilibrium outcome, we cannot disentangle pre-entry selection from post-entry workplace effects using variation in aggregate sorting alone.

This identification issue is not only a methodological problem. It is also fundamentally theoretical, because the selection versus influence question elides a basic sociological idea: that different people might end up in different occupations for different reasons. Some individuals may have the latitude to act on their political worldviews (to signal worth or claim expressive benefits), while others may face material constraints that make ideological congruence a secondary consideration at best. If this is correct, the self-selection process is itself class-stratified, and differential selection is a mechanism of occupational inequality: we cannot determine how politics affects occupational sorting without first identifying who treats their political worldview as a guide for major life decisions.

In this article, I take advantage of a unique empirical window to address both the methodological and theoretical problems at once. Instead of examining workers already sorted into occupations, I investigate *occupational preferences*—the vocational aspirations people form in early adulthood—to ask whether ideology shapes which occupations people aspire to pursue. As these preferences are

¹According to 2024 data from *Politics at Work*, which documents occupation-level estimates of two-party registration in the United States, almost 80% of people working as political scientists, social workers, or librarians are Democrats, while almost 70% of those working as clergy, oil operators, or commercial pilots are Republicans.

²Kitschelt and Rehm (2014) call these two pathways *weak* and *strong* theories of preference formation, respectively.

formed before people enter the labor market, any association between political ideology and career preferences cannot reflect post-entry workplace processes. Hence, I can examine whether political ideology shapes career aspirations, and whether this association is stratified along socioeconomic lines, investigating differential self-selection as a cultural mechanism of occupational sorting.

Using data on 9.5 million college freshmen entering U.S. institutions between 1976 and 2010, I study the occupational aspirations people articulate upon entering higher education. I show that political ideology, operationalized through students' self-identification as liberal, moderate, or conservative, systematically structures occupational aspirations. Using common benchmarks, I find that political stratification in 2010 is 23% as large as gender stratification and 79% as large as racial stratification. I further show that ideological sorting is not reducible to prior observable characteristics: compared to the counterfactual where political ideology is predicted by sex, race, socioeconomic background, religious upbringing, high-school achievement, and county of residence, ideological sorting retains 86% of its original magnitude. I map which occupations carry this political differentiation at large, and analyze whether the same set of occupational categories drives it across student cohorts.

Ideological sorting in occupational preferences, however, is much stronger among individuals from higher socioeconomic backgrounds. This implies that ideology functions differently depending on one's position in the stratification system: an expressive opportunity for some, a secondary concern for others. One interpretation of this income gradient is consistent with a luxury-good mechanism, where the salience of ideological congruence increases with material security (Enke, Polborn, and Wu 2025; Inglehart 1997). I show that, while liberals and conservatives differ along socioeconomic lines on the kind of occupational amenities they value, where high-income liberals prefer pro-social benefits while high-income conservatives prefer high earnings and prestige, such differences do not fully account for the income gradient. This suggests that ideology might be an expressive resource in the imagination of one's future, and that this expressive resource is unequally distributed.

These findings suggest that political ideology is a formative influence on career aspirations among college students and that this influence is structured by students' family background. While college freshmen are more affluent, liberal, and higher-achieving than their non-college peers (Bailey and Dynarski 2011; Scott 2022), this selection, if anything, makes for a conservative test: since material constraints are likely less binding in this group than among non-college youth, the socioeconomic gradient I document likely understates what would emerge in a population where such constraints vary more widely. Occupations pursued largely without a college degree nonetheless fall outside the design of this study, and whether these empirical results, and the gradient in particular, extend to occupational sorting among non-college youth cannot be resolved by the present analysis.

This article contributes to our understanding of political stratification in the labor market (Chinoy and Koenen 2024; Colonnelli, Neto, and Teso 2025), proposing that political culture structures how

young people come to envision the kind of work they want to do (Arold 2024; Doepke and Zilibotti 2008; Meriläinen and Mitrunen 2025). As such, I provide the first systematic evidence that ideology shapes occupational aspirations in early life. In doing so, I connect the status-attainment tradition (Blau and Duncan 1967; Haller and Portes 1973; Sewell, Haller, and Portes 1969)—long concerned with how social origins affect people’s occupational destinations—to contemporary work on political sorting, treating ideology as a cultural dimension of the attainment process. I further argue that the capacity to act on such commitments is itself unequally distributed, and political congruence is an expressive good available primarily to those whose material circumstances afford it (Inglehart 1997). These findings suggest that ideology is not merely a correlate of where people end up in the labor market, but a consequential force that shapes where they ultimately aspire to go.

Ideology as an Occupational Amenity

Occupations are central to the stratification of labor markets, simultaneously structuring earnings inequality (Weeden 2002), career trajectories (Spilerman 1977), and patterns of segregation across gender, race, ethnicity, credentials, and language (Alonso-Villar, Del Rio, and Gradin 2012; Hellerstein and Neumark 2008; Reskin, McBrier, and Kmec 1999). The allocation of individuals to specific occupational categories is thus a consequential stage in the production of labor market segregation, prior to workplace-level sorting within occupations. Accordingly, the mechanisms shaping this allocation are well-documented, emphasizing human capital (Becker [1964] 1994), status attainment (Blau and Duncan 1967; Haller and Portes 1973; Hout 1984, 2018), and social resources (Lin, Ensel, and Vaughn 1981), each structuring which occupations are *feasible* to pursue, as well as heterogeneity in individuals’ occupational aspirations based on personality fit (Holland 1959), non-pecuniary benefits (Rosen 1986), and preferences toward work and leisure (Doepke and Zilibotti 2008), each shaping which occupations are *desirable* to pursue, even for those facing similar constraints.

Understanding these allocation processes, however, requires us to examine the formation of particular occupational preferences—when people identify which careers to pursue in the first place—as these preferences shape subsequent sorting into distinct career tracks. Existing accounts typically emphasize how individuals evaluate occupations as bundles of attributes, giving more or less importance to compensation, intrinsic and experiential work qualities, and working conditions (Kalleberg 1977; Rosen 1986; York, Song, and Xie 2025), or focus on the match between personality traits and vocational domains (Holland 1959). Less explicit in these accounts, however, are the symbolic amenities of certain occupations: the perceived cultural character of the work itself, and the social affiliations and group identities occupations often signal. This is, in part, what Weeden and Grusky (2005) describe as the “cultural reputations” of occupations: lifestyle profiles or sociodemographic compositions that attract certain kinds of individuals to certain kinds of career preferences.

Such cultural reputations are, in many ways, organized around political ideology, since the lifestyle profiles or group identities occupations signal often become legible through their political valence—much like those “latte liberals” and “bird-hunting conservatives” (DellaPosta, Shi, and Macy 2015). Occupational choice may thus partly reflect a preference for ideological affinity, with actors gaining expressive payoffs from congruent occupational choices and incurring dissonance penalties from the incongruent ones (Akerlof and Kranton 2000; also see Schuessler 2000). Consider, for instance, the well-worn tropes of the liberal professor and the conservative cop: occupations whose political reputations reflect both the leanings of those who occupy them and the cultural meanings attached to the work itself. Such tropes persist, in part, because they are self-fulfilling: ideology shapes who is drawn to these occupations in the first place (Gross 2013). The question, then, is how pervasive this selection really is, and how political reputations shape the calculus of such decisions.³

Premise 1: Congruence as an Occupational Good

What makes one occupation politically preferable to another? We can examine this question with a simple model⁴ in which actors choose among occupations by trading off material attributes against symbolic attributes. In this model, occupations have bundles of “pecuniary” and “non-pecuniary” benefits (e.g., Rosen 1986) and actors use a latent utility index to choose among discrete alternatives (Bruch and Mare 2012; McFadden 1974). A person’s expected utility from pursuing an occupation, in this context, has two components: benefits from material attributes (e.g., earnings, occupational prestige, and working conditions) and benefits from symbolic attributes—chief among them what I call *political congruence* between one’s ideology and an occupation’s political reputation.

Definitions. By political ideology, I mean a person’s position on the liberal–conservative dimension, understood—following Hinich and Munger (1998)—as a cue or label that summarizes one’s issue positions. Ideology, in this sense, is a latent issue-predictive orientation over a left–right axis. By an occupation’s political reputation, I mean the political orientation widely perceived to characterize those who work in it (Gross 2013) and the political meanings ascribed to the occupation (Weeden and Grusky 2005): professors are thus typed as liberal and police officers are typed as conservative.

³The question of occupational attainment has long been the centerpiece of the stratification tradition (Blau and Duncan 1967; Featherman and Hauser 1978), some of which placed occupational aspirations at the heart of how social origins translate into destinations (Haller and Portes 1973; Sewell et al. 1969; but see Kerckhoff 1976). This tradition, however, was concerned almost overwhelmingly with the vertical consequences of aspiration formation (i.e., how much status and income a young person hoped to attain), leaving the cultural character of aspirations largely unexamined. I build upon this work and incorporate a reputational theory of occupational self-selection (Gross 2013) into the attainment model, which allows me to ask *not* how high young people aspire, but toward what kind of work, treating the capacity to act on one’s political ideology as a stratified dimension of the broader occupational attainment process.

⁴In Supplemental Materials A, I present a formal representation of this model using a rational choice framework, which provides one potential formalization of the following argument. The formalization fixes what ideological congruence is, how it enters occupational utility, and what it means for its salience to rise with one’s material security.

Combining the two, I define political congruence as the ideological proximity between a person's political ideology and an occupation's political reputation along this shared dimension.⁵

Congruence, so defined, functions as an occupational good. Occupations that align with one's politics offer expressive and identity-affirming benefits (Akerlof and Kranton 2000; Schuessler 2000), much as occupations that pay well or confer prestige offer material ones. In this sense, congruence is an occupational good among many other attributes, alongside earnings, prestige, autonomy, and security (Chinoy and Koenen 2024; Kalleberg 1977; Mas and Pallais 2017; Rosen 1986). Like many non-pecuniary rewards, it is something people trade off against pecuniary rewards. Therefore, the effect of congruence on aspirations depends on the *weight* it receives *relative to the others*.

These considerations lead to the first premise for a theory of ideological self-selection:

Premise 1: Ideological congruence is one good among the many an occupation offers; its influence on aspirations depends on the weight it carries relative to alternative rewards.

Premise 2: Congruence as a Luxury Good

The relative weight assigned to political congruence compared to material considerations, however, need not be constant across individuals, as material and expressive goods speak to different needs. Concerns about earnings, security, or stability are ultimately about the material conditions of one's future, whereas ideological congruence allows occupational choice to express and affirm a person's values (Inglehart 1997; Schuessler 2000). The relative importance of these goods should therefore depend on the material circumstances under which occupational aspirations are formed.

Why would material circumstances alter the relative balance between these goods, though? As the material conditions of the future become more secure, consequential choices can be oriented more strongly toward self-expression, identity, and the realization of one's values (Inglehart 1997). Put differently, material security may alter *what people seek from consequential life choices*. Applied to the context of occupational aspirations, this means that ideological congruence should become a more salient basis for evaluating occupations among students from more advantaged backgrounds. As formalized in Enke et al. (2025), this is a luxury-good effect: political considerations carry greater weight as material security rises, without presuming stronger ideological commitments.

These considerations lead to the second premise for a theory of ideological self-selection:

Premise 2: The weight congruence carries rises with material security, making it a luxury good in the formation of occupational aspirations.

⁵Such an argument for "congruence" assumes a symmetric loss in utility when people assess whether the occupation is to their "left" or to their "right." While asymmetric preferences might exist—e.g., conservatives avoiding left-coded occupations more so than liberals avoid right-coded ones—the expectation holds under either specification.

Empirical Implications for Self-Selection

Let J_{it} denote an actor i 's intended occupation at time t , where ω represents the set of occupational categories available in the relevant choice set. I use occupational *intentions* in early adulthood as a measure of occupational *aspirations*. Under this interpretation, when an actor reports occupation j as their intended career from a list of occupational categories presented in ω , this reveals a relative evaluation over occupations, as defined via the actor's latent utility.⁶ This process at the micro-level produces an observable implication at the macro-level, which I state as the first hypothesis:

Hypothesis 1: Ideological positions systematically structure occupational aspirations.

Of course, because ideological congruence is only one component of utility, its practical significance depends on what else is at stake in the general choice setting. As I proposed before, ideology may be more or less important depending on one's income position. If this is correct, individuals from low-income families should systematically select into occupations on non-ideological grounds, making differential self-selection a mechanism of occupational stratification and ideology a class-stratified cultural resource. These considerations mean that, at the macro-level, we should observe:

Hypothesis 2: The systematic association between ideology and occupational aspirations is stronger among individuals from higher-income families.

Alternative Interpretations. The income gradient in Hypothesis 2 invites a taste reading, in which higher-income individuals assign greater weight to congruence. Yet the gradient alone admits two alternative interpretations—each resting on a different individual mechanism. While these cannot be ruled out entirely, I take each interpretation in turn and investigate alternative predictions.

The first reading is *diminishing marginal utility*. Higher-income students may sort more strongly not because they weight ideological congruence more, but because the marginal value of an additional dollar is lower for them, so material returns weigh less in their choices (people from high-income families, for instance, might forgo material benefits, simply because they expect more intergenerational wealth transfers).⁷ On the strength of the gradient alone, the two are thus indistinguishable. That said, there is a weaker version of this interpretation that I can investigate: if people sort more *only because* they assign less weight to material rewards, this re-weighting should be general rather than specific to political ideology: as material concerns weaken, *every* non-pecuniary consideration

⁶While occupational intentions will incorporate anticipatory constraints—e.g., individuals avoid stating intentions for careers they perceive as entirely infeasible—they still reveal relative preferences within one's perceived choice set.

⁷However, this does not make the diminishing marginal utility argument a more parsimonious default. As Enke et al. (2025) convincingly argue, the concavity of utility does not, by itself, result in the income gradient: while it lowers the marginal value of an additional dollar, the absolute stakes of a career decision are still larger for better-resourced individuals, and the two effects need not cancel. Arguing that the weight on congruence rises with material security is thus no less principled an assumption than the assumption of concavity overwhelming the absolute income effect.

should gain prominence. An income gradient that is instead markedly stronger for ideology would rule out this deflationary version of diminishing marginal utility, leading to the hypothesis:

Hypothesis 3: The income gradient in ideological sorting is steeper than the income gradients in sorting along other non-pecuniary occupational valuations.

The second reading is *differential perception*. Congruence, according to this account, can guide aspirations only if occupations are perceived with clear ideological signals in the first place (Gross 2013; Weeden and Grusky 2005); higher-income students may then sort more strongly not because they weight congruence more, but because their reputational map differs from that of their low-income peers.⁸ Put differently, students from different backgrounds may perceive political reputations in different ways. On the strength of the gradient alone, taste and perception are again indistinguishable. This reputational difference, however, can leave one behavioral signature that I may detect: if higher-income students act on a differently structured map, the association between ideology and aspirations should differ in *structure* across income groups, with different occupations carrying the sorting at different income levels. If students instead act on a shared political map, the association should retain a common structure, with income differences confined to its overall strength⁹:

Hypothesis 4: Income differences in sorting reflect the strength of a shared ideology and occupation association, rather than differences in the structure of that association.

The Empirical Setting

My empirical approach centers on *the occupational aspirations of college freshmen* as they enter higher education. This is a valuable population for two reasons. First, college students are systematically encouraged to formulate their occupational intentions: they must declare majors, seek internships, and make consequential human capital investments—all requiring them to articulate occupational goals. This institutional pipeline makes occupational aspirations a particularly meaningful indicator of how young people’s life-course trajectories may unfold. Second, these career aspirations are elicited *before* the labor-market entry simultaneously for most students, contrary to more heterogeneous trajectories among the non-college youth. This, in turn, presents a unique empirical window for observing preference formation prior to workplace socialization or firm-level sorting.

⁸Such perceptions might form through a variety of mechanisms, such as personal observations of who works in a given occupation, media representations, ideological transmission from political actors, or politicization of certain careers via public opinion campaigns. For instance, Gross (2013) provides a detailed historical documentation of how right-wing political operatives engaged in activities that politicize the perceptions of college faculty as distinctively liberal.

⁹This test cannot rule out an alternative form of perceptual heterogeneity in which students share the same reputational map, but higher-income students perceive its ideological signals more clearly (Lynn and Ellerbach 2017; Lynn, Shi, and Kiley 2025). Such a mechanism would preserve the occupational structure of sorting while increasing its overall strength, mimicking a luxury-good prediction. I return to this possibility in more detail in the concluding sections.

Of course, these analytical advantages come with one important limitation. College students are a selected population: they are more affluent (Bailey and Dynarski 2011), more liberal (Scott 2022), and presumably more amenable to expressive choice than their non-college peers. The question of whether and how the findings documented below generalize to non-college youth is therefore an empirical one. If anything, however, this selection makes the central argument, that political sorting in aspirations is stronger among higher-income students, analytically more conservative: because the sample is concentrated in a population for whom material constraints are probably less binding than among the non-college youth, the income gradient documented below may understate what would presumably emerge in a sample with greater material variations and constraints.

The Survey

I use survey data from *The Freshman Survey* (TFS), administered by the Higher Education Research Institute (HERI) at the University of California, Los Angeles. TFS is designed to collect information about the population of first-time, first-year students at U.S. institutions, with almost 2,000 colleges and universities participating at various points across the study period. Importantly, these surveys are conducted before students begin their classes in their first year, so they capture students' views and preferences before they are fully exposed to the college environment. The final sample includes observations from 1976 to 2010, with almost 9.5 million participants aged 16 to 20.¹⁰ TFS provides sampling weights that correct for institutional attrition and coverage, as participation is voluntary at the college level and response rates vary over years and across institutions. I present unweighted estimates to retain all the available observations, while also reproducing all findings with weighted estimates.¹¹ Supplemental Materials B provides relevant details and descriptive statistics.

TFS has recently been used in several papers examining political ideology among college students. Acton, Cook, and Ugalde Araya (2025) study how political ideology shapes college choices, finding increasing political differentiation among college attendees. Goldstein and Kolerman (2025) focus on the effects of academic fields on political attitudes, finding differential changes in political views during college across different majors. More closely related to the present article, Reny et al. (2025) find that individuals intending a career as a police officer report more right-wing political attitudes than their peers. In this article, I draw from this line of research and provide a systematic theoretical investigation of why, and for whom, political ideology informs occupational aspirations.

¹⁰I made three important decisions while processing the data. First, I removed survey years before 1976, as the survey item used to elicit participants' occupational preferences had substantial revisions in earlier years. Second, I dropped people aged 21 and older ($\approx 5.7\%$ of the initial sample), so that I can investigate early adults as my target population. Finally, I dropped individuals with missing responses on political ideology or occupational aspirations ($\approx 12\%$).

¹¹Note that about 30% of observations have missing weights, which means that weighting the data reduces the effective sample size quite significantly. This reduction mainly results from excluding colleges without sufficient coverage.

Measurement

Occupational Aspirations. To capture aspirations, I use a survey question asking freshmen about their “intended careers.” This item captures 48 unique occupational categories over time, provided as pre-aggregated units in the TFS data. Since response options and category labels in the original questionnaires change across survey waves—for instance, “Accountant or Actuary” in earlier waves becomes simply “Accountant” in later ones—I prepared an LLM-assisted protocol to document the original questionnaire wording and a response crosswalk, which I provide in the reproduction files. Ultimately, I focus on the 41 categories that appear consistently over years, which account for more than 98% of category responses, and reassign the remaining ones into broader categories.¹²

There are two sources of measurement uncertainty about this strategy. First, since original response options change across survey waves, it is likely that large breaks in the time series reflect instrument redesign rather than genuine preference change. Looking at the marginals of occupations over time, the most consequential such period spans 1995 to 2000 (see Figure B1 for composition across years). I therefore flag this time period as potential instrument drift in analyses below. Second, the “Other” and “Undecided” categories are aggregate buckets, with their composition shifting in response to the available response options over time. I thus drop undecided responses from most analyses and verify that the findings are not sensitive to the inclusion of the “Other” response category.

Political Ideology. To capture ideology, I use an item asking students to characterize their political ideology on a five-point scale: *far left*, *liberal*, *middle-of-the-road*, *conservative*, or *far right*. Because the extremes of this scale feature relatively few respondents (3.9% of the complete sample), I collapse it into three main categories: liberal (combining far left and liberal), moderate (middle-of-the-road), and conservative (combining conservative and far right). Figure B2 presents the marginals of these categories over time. Of course, it is likely that the *meaning* of these categories changed across years: someone identifying as “liberal” in 1980 may not be equivalent to someone doing so in 2005. I thus complement the main analyses with 12 issue-specific questions—ranging from capital punishment to the legality of same-sex relations—to measure issue-specific pathways. Since these items appear in different survey years, I cannot construct a composite scale, so I replicate the main analyses using each item separately. In Supplemental Materials C, I provide the complete list of these issue items, and the specific question wording and response options used to measure these positions.

Socioeconomic Background. I measure socioeconomic position using parental household income, as reported by students upon entering college. Since parental income is the primary determinant of

¹²I recoded *college administrator* into *other*, *federal/state/local government official* into *other*, *interpreter* into *other*, *laborer* into *other*, *physician assistant* into *medical/dental assistant*, *semi-skilled worker* into *other*, and *statistician* into *research scientist*. Dropping observations with these responses from the analyses below does not change the substantive conclusions of the study. Table B2 in the Supplemental Materials B presents the marginal distribution of occupational intentions.

material circumstances for young adults, I use it as the main item for socioeconomic background.¹³ TFS measures parental income in bracketed categories (e.g., “\$20,000–\$24,999”), and these bracket definitions change across survey years. To ensure comparability, I first assign each income category its midpoint¹⁴ and convert these values to constant 2010 dollars with the Consumer Price Index. I then transform parental income into a within-year decile rank, following the income-rank approach common in stratification research (Chetty et al. 2014): for each period, I compute fractional income ranks and assign students to deciles or quartiles based on those ranks. This transformation allows me to account for inflation and bracket redesigns over time, ensuring that a student, for instance, in the 8th decile in 1980 and another one in the 8th decile in 2005 are comparably positioned relative to their peers in the same year.¹⁵ While within-year income rank is consistently used throughout the analyses below, I also provide supplemental analyses using real parental income (reported income in constant 2010 dollars adjusted with the U.S. Bureau of Labor Statistics’ CPI estimates).

This measure is necessarily imperfect: students may not know their parents’ exact income, though reporting error is likely greater for finely grained dollar values than for broad economic positions. I rely primarily on within-year income rank rather than absolute income values for this reason, as this measure treats reported income as an *ordinal* indicator of background, which is less demanding than assuming precise knowledge of family income. Moreover, since the argument concerns how students evaluate future careers under *perceived material constraints*, students’ own perceptions are arguably the more relevant theoretical quantity than their parents’ precise income levels.

Analytic Strategy

To evaluate the hypotheses within a common framework, I define *the degree of political differentiation in occupational aspirations* as my empirical estimand: if ideological congruence carries weight in the formation of aspirations, it implies that the distribution of occupational aspirations should systematically vary over ideological positions. I quantify this differentiation using mutual information M , which I calculate from the contingency table of occupational categories J and political positions P . More intuitively, M captures the degree to which the joint distribution of aspirations and ideology departs from independence. Because M may be computed within any partition of the data—survey years or population subgroups—it provides a unifying framework for the analyses that follow.

¹³Students respond to the item: “What is your best estimate of your parents’/guardians’ total income last year? Consider income from all sources before taxes.”

¹⁴For open-ended response options (e.g., “\$150,000 or more”), I assign plausible midpoint values. Because the primary analyses rely on within-year income ranks rather than these dollar values themselves, the results are not sensitive to these choices. For alternative analyses that use absolute values, I discuss reasonable midpoint alternatives below.

¹⁵Since income is measured in grouped bins rather than continuous values, respondents within the same bin received the same decile assignment. One comparability problem is that the top income categories are coarser in early survey years, meaning that the 10th decile is more heterogeneous in early years than in later ones. This likely makes tests of Hypothesis 2 much more conservative in the early part of the study window, compared to ones in later windows.

Let n_{jp} denote the number of individuals with occupational preference $J = j$ and political ideology $P = p$, and let N be the total count of individuals. These quantities allow me to calculate $q_{jp} = n_{jp}/N$ (the joint share), $q_j = \sum_p q_{jp}$ (the marginal share for occupation j), and $q_p = \sum_j q_{jp}$ (the marginal share for ideology p). With these values at hand, the mutual information M is then defined as

$$M(J;P) = \sum_j \sum_p q_{jp} \log \frac{q_{jp}}{q_j q_p} \quad (1)$$

Mutual information has several properties that make it well-suited for this study. First, it imposes minimal structure on the data: both occupational aspirations and political ideology are categorical variables, and M requires no assumptions about how occupations are scaled and ranked, capturing any systematic dependence between the two. Second, M has strong credentials in the measurement of sorting due to its *decomposability* (see [Mora and Ruiz-Castillo 2011](#); [Reardon and Firebaugh 2002](#)). That is, the aggregate value may be additively decomposed into occupation-specific contributions, identifying which occupational categories carry the associations. This allows me to move between the overall magnitude of political differentiation and its composition across empirical settings.

These properties notwithstanding, M is not a normalized measure: its magnitude depends on the marginal distributions of both variables. I address this limitation in two ways. First, in supplemental analyses, I complement these raw values with normalized variants that divide M by the entropy of the occupational distribution or by the entropy of ideology groups, bounding the measure and adjusting for differences in marginal structure. Second, I empirically anchor interpretation through benchmarking, where I calculate M between aspirations and gender, as well as between aspirations and race, using the same samples, the same occupational categories, and the same procedure.

Two final notes concern *inference* and *weighting*. Since the analytic sample features millions of observations, sampling uncertainty is negligible for most aggregate quantities; the more consequential concern for the analyses is that mutual information is biased upward in finite samples, particularly when contingency tables are sparse. I thus verify that the *expected value of M* under independence is effectively zero in survey years to ensure that the observed values do not reflect finite-sample noise. I report percentile confidence intervals from bootstrap resampling, although in certain cases these are too small to be visible. Finally, as noted before, TFS sampling weights are missing for roughly 30% of the observations. I therefore present unweighted calculations as my baseline quantity.

Mapping Ideological Sorting

I start with Figure 1, which presents the occupation-level ideological differentiation in TFS, pooled across the 2005–2010 survey waves. The figure plots, for each occupation, the share of conservative




Figure 1: Ideological Sorting in Occupational Aspirations

Notes: The figure presents the distribution of ideological sorting in occupational aspirations, computed from the pooled 2005–2010 data using share conservative minus share liberal as the main statistic. Values above zero indicate occupations with higher conservative than liberal share, values below zero indicate the reverse pattern, and the dashed vertical line represents equal shares of liberals and conservatives. Point size is proportional to occupation size.

students minus the share of liberal students, with group shares computed over everyone, including moderates. The resulting score ranges from –100% (all liberal) to 100% (all conservative): careers to the right of the reference line are those where conservatives outnumber liberals, while categories to the left reflect the opposite ordering, with point size being proportional to occupation size.

Figure 1 suggests that there is substantial ideological variation in occupational aspirations. On the right, we see military, clergy, homemaker, farmer or forester, and members of protective services as the most distinctive occupations. On the left, we see foreign service workers, artists, college faculty, as well as writers and social workers. In contrast, several occupational categories cluster fairly close to zero, which indicates that the share identifying liberal and the share identifying conservative are roughly similar for these cases. Notably, several high-prevalence occupational categories—such as engineers or business executives and managers—indicate almost zero differentiation.¹⁶ The overall finding, then, is one of substantial though uneven differentiation in occupational aspirations.

While Figure 1 reveals where political differentiation is concentrated, it does not speak to the overall magnitude of ideological sorting. To address this question, I now return to mutual information M . Figure 2 presents the level of political sorting in occupational aspirations over time. The top panels vary how moderates are incorporated into the analyses: Panel A uses all three ideology categories, Panel B restricts the TFS sample to liberals and conservatives only, and Panel C contrasts moderates against ideologues (liberals and conservatives combined). The bottom panels provide benchmarks against sorting along occupational sorting by gender and occupational sorting by race. I shade the 1995–2000 period to emphasize the substantial changes in the instrument in this time period.

Figure 2 suggests that the overall magnitude of political sorting is strong and consequential. To anchor this claim, consider the 2010 estimates of mutual information alongside common benchmarks. When all three ideology groups are included, the magnitude of ideological sorting is roughly 23% of the magnitude of gender sorting—by far the largest axis of occupational segregation in general. Against racial sorting, the magnitude is even more striking: political sorting is 79% as large as racial sorting. Restricting the focus to liberal and conservative students only amplifies these figures considerably, with political sorting rising to 32% as large as gender sorting and 111% as large as racial sorting. These results highlight that political ideology is a strong axis of occupational sorting.

How do these patterns change over the years? I find that the temporal trajectories depend critically on how moderates are treated. When all three ideology groups are included (Panel A), M remains virtually flat from the mid-1970s through 2010. Excluding moderates (Panel B), however, reveals a starkly different picture: political sorting among liberals and conservatives nearly doubles over the same window, suggesting that ideological polarization in occupational aspirations is concentrated in the liberal–conservative contrast. Consistent with this interpretation, Panel C shows that sorting between moderates and ideologues decreases by more than 50% over time, suggesting that political sorting among students relies primarily on liberals and conservatives diverging across years.¹⁷

¹⁶While Figure 1 provides unweighted estimates, the weighted ones are identical, with Spearman's $\rho > 0.98$. Similarly, excluding moderates from the denominator leads to similar findings (with Spearman's $\rho > 0.99$).

¹⁷I reproduce the analyses using twelve substantive issue positions in place of ideological identification in Supplemental Materials C, showing that policy views are similarly informative of occupational aspirations. A share of the increase I document in Figure 2, however, should be interpreted with caution: the *alignment* between ideological identity and

Ideology




Figure 2: Temporal Trajectories in Ideological Sorting

Notes: The figure presents changes in ideological sorting over time, and benchmarks ideological sorting against sorting by gender and sorting by race. Each panel represents yearly M values between occupational aspirations and the relevant grouping variable (described on the bottom-right corner of each panel), as weighted and unweighted series. Participants who are “undecided” in their career aspirations are excluded. The shaded band marks 1995–2000 to highlight temporary changes in the occupational preference instrument in the surveys. “Ideologues” in Panel C refer to an alternative coding where liberals and conservatives are combined into one group for comparisons with moderates.

Robustness Checks. In Supplemental Materials C, I address the concern that the trajectories of M reflect shifting marginal distributions rather than changes in the occupation–ideology relationship. I show that the effective number of occupational categories remains stable across the study window, normalized variants of M reproduce the temporal trajectories, and a decomposition of changes in M demonstrates that the divergence between liberals and conservatives reflects structural change in

substantive issue positions itself strengthened over this period, meaning that part of the observed trend likely reflects better sorting of students into ideological labels rather than the politicization of occupational aspirations as such. In any case, issue-level sorting presented in Figure C10 shows that political beliefs are consequential for sorting.

the occupation–ideology association rather than movements in the marginals. Across these checks, I find that the substantive conclusions supporting Hypothesis 1 remain virtually unchanged.

Ideological Sorting and Sociodemographic Background

I now turn to the question of what accounts for political differentiation in occupational aspirations. It might be the case that ideological sorting is largely a reflection of pre-existing sociodemographic differences, and much of political sorting is attributable to these compositional differences. I study this question using a *decomposition* method for M . Specifically, I decompose M into two components: the share of occupation–ideology association attributable to observable predictors of ideology and the share that persists net of these factors (following Åslund and Nordström Skans 2009).

In doing so, I estimate the propensity to identify as liberal, moderate, or conservative given a set of factors, $\Pr(P = p \mid X)$. I then construct a counterfactual occupation-by-ideology table using model predictions, a strategy that holds the distributions of occupations fixed, but replaces the observed ideology with the predicted ideology.¹⁸ The counterfactual M_C computed from this table captures the ideology and occupation dependence implied purely by antecedent empirical attributes X . The end quantity is what I call residual mutual information, or M_R , which is represented as:

$$M_R = M - M_C \quad (2)$$

This approach is best understood as a descriptive decomposition rather than a causal identification strategy. M_R captures the ideological sorting that persists after removing the pathways associated with X , allowing me to examine how much of the aggregate association covaries with X .

Figure 3 presents the main results in the context of sociodemographic characteristics. In each row, I sequentially expand the predictor set X : (1) sex interacted with race, (2) parental socioeconomic status (measured through fixed effects for household income decile, educational attainment of the mother and the father, as well as their reported occupations), (3) religious upbringing (measured through parental religious affiliation), (4) achievement in high school, measured through students' self-reported high-school GPA, and (5) county of residence.¹⁹ The figure suggests that these characteristics explain little: the residual mutual information remains close to the baseline in both panels, implying that ideological sorting is not an artifact of prior sociodemographic differences.²⁰

¹⁸More precisely, I use pairwise logits for predicting liberal, moderate, and conservative identification (then combining these estimates via softmax, which allows me to avoid a computationally intensive multinomial logit), while using a simple logistic regression predicting the propensity to identify as liberal for the two-category analyses.

¹⁹I assigned all missing observations to an explicit “missing” category for estimation.

²⁰In Supplemental Materials C, I present the weighted decomposition analysis, with substantively similar conclusions.




Figure 3: Decomposition of Ideological Sorting by Sociodemographic Characteristics

Notes: The figure presents a decomposition of ideological sorting in occupational preferences by sociodemographic background characteristics. Baseline refers to observed M , while other rows present residual M after sequentially expanding the set of predictors used to construct counterfactual ideology distributions. Panel A uses all three ideology categories, while Panel B restricts the sample to liberals and conservatives. Missing values are included in the estimations.

The Occupational Structure of Ideological Sorting

The analyses so far establish that sorting is sizable and increasingly organized around the liberal-conservative axis. That said, M remains an aggregate quantity: it summarizes the association without showing which occupations produce it. One natural question to ask is thus whether this sorting is something diffuse, spread thinly across the occupational structure, or concentrated in a small set of politically typed careers. To answer this question, I exploit the decomposability of M (Mora and Ruiz-Castillo 2011) by rewriting M as a weighted sum of occupation-specific contributions:

$$M(J;P) = \sum_j q_j L_j, \text{ where } L_j = \sum_p \frac{q_{jp}}{q_j} \log \frac{q_{jp}/q_j}{q_p} \quad (3)$$

where L_j captures how much occupation j 's ideological composition departs from the population's ideological distribution and q_j weights this departure by the occupation's size. Occupation j 's share of overall sorting is then $q_j L_j / M$, with shares summing to one. An occupation might therefore carry a large share of M in two ways: by being extremely "typed" (high L_j) or by being popular enough that even very modest differentiation aggregates into a substantial contribution (i.e., high q_j).

I calculate these shares within five-year windows, which Figure 4 presents as a heatmap. The figure shows that ideological differentiation is heavily concentrated, with 25% of all occupations carrying roughly 70% of sorting. More importantly, while some of these occupations are recognizable from Figure 1 (e.g., the military or actors), the decomposition also reveals contributors to political sorting




Figure 4: Occupation-Specific Contributions to Ideological Sorting

Notes: The figure presents the decomposition of ideological sorting into occupation-specific contributions. Each cell represents an occupation's share of the aggregate mutual information M in the corresponding period, with shares summing to 100% within each period and darker cells indicating larger contributions. Participants who are "undecided" in their career aspirations, as well as those in the residual "Other" category, are excluded from the decomposition.

that raw differentiation obscures: large categories such as lawyers carry substantial shares through their sheer popularity. I also find that the same set of occupations anchors the association over time, though their relative weights change: for instance, the contribution of legal careers declines steadily across cohorts, while the military rises to become the single largest carrier by the late 2000s.²¹

Are these aspirations consequential for the labor market? In Supplemental Materials C, I compare occupation-level ideological differentiation in aspirations with the realized partisan compositions of the same careers, estimated from voter-registration records linked to employment data. I find a cross-occupation correlation of $\rho \approx 0.88$: in other words, the occupations that attract conservative-identifying students are largely the careers staffed by registered Republicans, and likewise for their liberal counterparts. Aspirations, therefore, closely anticipate the ideological map of the American

²¹In Supplemental Materials C, I present the weighted decomposition analysis, with substantively similar conclusions.

occupational structure. While this does not establish the aspiration–destination match, it provides evidence consistent with it. Taken together, these findings lend strong support to Hypothesis 1.

Ideology as a Luxury Good

I now turn to the question of whether ideological sorting in occupational aspirations varies across socioeconomic background, testing the existence of an “income gradient” in ideological sorting.

In Hypothesis 2, I predict that political ideology and occupational aspirations covary more strongly among individuals from higher socioeconomic backgrounds. To explore this prediction, I compute mutual information M between ideology and aspirations within each parental income decile, using within-year parental income rank as the measure of socioeconomic position. If the relative salience of political congruence increases with socioeconomic background, mutual information M between political ideology and aspirations must increase monotonically as a function of parental income.

Figure 5 presents the findings, using weighted and unweighted samples across columns, and varying ideology groups across rows. These results provide strong support for Hypothesis 2: students from the highest-income families present roughly two to three times the level of ideological sorting in occupational aspirations compared to students from the lowest-income families. Once we focus on the difference between the top and bottom rows, we observe important differences, as excluding moderates both raises the overall sorting in M , and steepens the income gradient substantively.

Robustness Checks. In Supplemental Materials D, I probe these results further. The first concern is mechanical: the breadth of students’ occupational choice sets might move M across deciles without any change in the weight of ideology.²² Normalizing M by within-decile political ideology entropy or by within-decile career entropy preserves this gradient, ruling out these compositional alternatives. The second concern is pooling: because the analyses combine all survey waves, the gradient could reflect differential trends across time. Estimating the gradient within five-year periods, however, shows that it is present in every time period. The third concern is measurement: it might be that using income rank obscures real income differences. Replacing within-year decile ranks with CPI-adjusted real parental income reproduces the gradient. Finally, I show that these patterns are not an artifact of sample composition: re-estimating it on random 30-occupation subsets yields the same trajectory, and while the propensity to report an “undecided” career varies across both ideology and income rank, having these students in the analyses does not change the gradient. Across these specifications, the income gradient holds, providing strong support for Hypothesis 2.

²²This is an important possibility to adjudicate: higher-SES students may simply aspire to a wider range of occupations. If the gradient survives normalization by within-decile occupational entropy, that would give strong evidence against this composition-based story




Figure 5: Ideological Sorting Across Socioeconomic Background

Notes: The figure presents mutual information M between political ideology and occupational aspirations across household income, which I calculate as within-year income rank across 1976-2010. The estimates include bootstrap confidence intervals from 1,000 draws, as well as a linear fit to describe the income gradient. Panels A and B: all three ideology categories. Panels C and D: liberals and conservatives only.

Extensions. I extend these gradient analyses in three directions in Supplemental Materials D. First, because parental income and parental education are strongly correlated, one may ask whether the gradient reflects material security or the parental educational background. Stratifying individuals by parents' college attainment shows that M still rises with income within both education groups, though, as expected, sorting is substantially higher among students from college-educated families. Second, I demonstrate that high-income students are far more likely to aspire to their parents' own occupations: roughly 12% of students report a parental occupation as their intended career, a share rising from about 6% in the bottom income quartile to about 21% in the top. The gradient, however, is not an artifact of this reproduction, as excluding these students leaves the gradient intact. Third, the finding is not specific to ideological labels: replacing ideological identification with the twelve issue positions results in a substantial income gradient for nearly every belief and preference, with ideological identification itself presenting, as expected, the steepest income gradient.

Income Gradient for Gender and Race

One might worry that the income gradient reflects something *generic* about students from affluent families rather than anything *specific* to political congruence: higher-income students might simply hold more differentiated occupational preferences across the board, in which case sorting along any dimension would rise with parental income. To evaluate this possibility, I repeat the previous exercise for the two benchmark axes of occupational differentiation, computing M between aspirations and gender, as well as between aspirations and race, within each parental income decile.

Figure 6 demonstrates that neither axis behaves anything like political ideology: sorting by gender follows an inverted-U shape across the income distribution, while sorting by race declines steeply and monotonically: aspirations in the top decile are sorted by race at roughly 40% of the magnitude observed in the bottom decile. These results imply that the income gradient is *not* a general feature of occupational differentiation: put differently, among the major axes of aspiration sorting such as gender and race, the income gradient seems to be highly specific to ideological identification.²³

Alternative Occupational Valuations

While Figure 5 presents strong evidence that ideological sorting rises with household income, that pattern, in itself, cannot establish that ideological congruence becomes more important as material security increases. A simpler explanation is that high-income liberals and conservatives differ more strongly than low-income ones in the kinds of occupational attributes they value. If this is indeed the case, the income gradient could partly reflect differences in such valuations, which show up as ideological differences. To properly address this issue, we need to ask whether the income gradient persists after accounting for differences in how students evaluate occupational attributes.

To do so, I draw on data from two survey years (1989 and 1990). In these two years, students rated thirteen career attributes—including material concerns such as high anticipated earnings or career advancement, as well as intrinsic dimensions such as the ability to work with ideas and making an important contribution to society—from *not important* to *essential*, providing importance judgments about each. I regress each standardized attribute rating on ideology, household income rank that I 0–1 normalized, and their interaction, adjusting for year, sex, race, and age fixed effects. This gives, for each attribute, the ideology gap at the bottom versus the top of the income distribution. Table 1 presents these findings, showing several important patterns. We see that the conservative-liberal gap in material attributes is modest at low income but expands substantially at the top. Similarly, liberals consistently place greater weight on intrinsic and altruistic attributes, and this gap widens

²³A compositional concern applies to these benchmarks as well: within-decile diversity in gender and race mechanically bounds M , and high-income deciles are, for instance, less racially diverse. In Supplemental Materials D, I normalize each axis by its within-decile group entropy, showing that the same patterns remain after normalization.




Figure 6: Gender and Racial Sorting Across Socioeconomic Background

Notes: The figure presents mutual information M between gender and occupational aspirations (top panels), and race and occupational aspirations (bottom panels), across household income, which I calculate as within-year income rank across 1976-2010. The columns present unweighted and weighted estimates. The estimates include bootstrap confidence intervals from 1,000 draws, as well as a quadratic fit to describe the income gradient.

with parental income. These patterns raise the possibility that part of the observed income gradient in sorting reflects these differences rather than the relative salience of political ideology.

Figure 7 tests whether the income gradient is primarily an artifact of such differential gaps between liberals and conservatives. Using the 1989–1990 sample, I replicate the residual mutual information exercise I introduced before within each parental income decile²⁴, this time accounting for the full set of thirteen career-attribute ratings. If the income gradient primarily reflects these compositional differences, then accounting for them should substantially flatten the income profile over M . Figure 7 shows that this is not the case. Conditioning on the attributes reduces the overall level of sorting, but the income gradient is almost entirely preserved. In other words, while occupational-attribute valuations may explain a portion of the observed ideological sorting, they do not fully explain the gradient, providing further evidence for the differential effects of political congruence.

²⁴Of course, the attribute ratings are themselves endogenous to ideology, so this represents a conservative test.

Table 1: Ideological Gap in Occupational Attributes Across Low and High Income

Attribute	Gap at Low Income	Gap at High Income
Available Job Openings	0.11	0.18
Rapid Career Advancement Possible	0.07	0.34
High Anticipated Earnings	0.06	0.47
Well-Respected Prestigious Occupations	0.07	0.41
Great Deal of Independence	−0.23	−0.16
Chance for Steady Progress	−0.04	0.12
Can Make Important Contribution to Society	−0.10	−0.24
Can Avoid Pressure	−0.02	−0.09
Can Work with Ideas	−0.21	−0.32
Can Help Others	−0.05	−0.18
Able to Work with People	−0.03	−0.04
Intrinsic Interest in the Field	−0.08	−0.16
The Work Would Be Challenging	−0.07	−0.09

Notes: The table presents the conservative–liberal gap in evaluations at the bottom (income rank = 0) and top (income rank = 1) of the household income distribution, calculated from a regression where attribute evaluations are regressed on ideology, household income rank, and their interaction, adjusting for year, sex x race, and age. In each estimation, the outcome variables are standardized to have a mean of 0 and a SD of 1. Positive values imply conservatives rate the attribute higher than liberals, and negative values indicate the reverse. Data from 1989–1990. N = 406,104.

Investigating Alternative Mechanisms

In this section, I turn to the question of alternative mechanisms, investigating Hypotheses 3 and 4 to evaluate the respective roles of diminishing marginal utility and differential perception.

The Specificity in Ideological Sorting

While the analyses so far test whether differential valuations confound the gradient, the diminishing marginal utility interpretation introduced earlier claims a broader prediction that these analyses do not yet test: if higher-income students sort more on ideology *only because* material rewards weigh less in their choices, this re-weighting should be general rather than specific to politics: sorting along every non-pecuniary consideration should rise with parental income. In Hypothesis 3, I stated an alternative: that income gradient should be markedly steeper for ideology than for other non-pecuniary attributes. To test this prediction, I estimate the income gradient separately for each dimension: for political ideology and each of the eight non-pecuniary attribute valuations, I compute normalized M^{25} between that dimension and occupational aspirations within each household

²⁵Because the ideology measure has three categories while the attribute ratings have four, raw M values are not directly comparable across dimensions. I therefore normalize each dimension's sorting by its own entropy before computing




Figure 7: The Income Gradient in Ideological Sorting

Notes: The figure presents observed and residual mutual information between ideology and aspirations across parental income deciles in the 1989–1990 sample. Residual M is computed using the same decomposition strategy described in earlier sections, now applied within each income decile using the thirteen standardized career-attribute ratings. Panel A uses all three ideology categories, while Panel B restricts the sample to liberals and conservatives.

income decile, and summarize the final gradient as the linear slope of M across deciles. I bootstrap this entire estimation pipeline to obtain confidence intervals for the final income gradients.

Figure 8 presents these findings, which are inconsistent with a general weighting of non-pecuniary considerations. Under the deflationary reading, in which the income gradient reflects nothing but the declining marginal value of material rewards, no dimension should stand out and none should decline. Instead, we see that the income gradient for ideology is roughly three times the size of the steepest attribute, that most attribute gradients are closer to zero, and that the gradient for several considerations is either non-significant or negative. The result that ideology has a distinct character therefore rules out this deflationary alternative, providing direct support for Hypothesis 3.

This finding does not, of course, rule out diminishing marginal utility as such. The analysis instead is directed at a deliberately deflationary reading of the diminishing marginal utility hypothesis: the claim that the income gradient in ideological sorting is merely one observable instance of a general tendency for affluent students to sort more strongly on all non-pecuniary rewards. The results are inconsistent with that account, because ideology shows a markedly steeper gradient than the other occupational valuations. That said, one could still argue that material security lowers the marginal

slopes, placing all gradients on a common scale, which allows me to make direct quantitative comparisons.




Figure 8: Income Gradients in Sorting Across Non-Pecuniary Valuations

Notes: The figure presents the income gradient in sorting for ideology and for each of the eight non-pecuniary attribute ratings, using the 1989–1990 sample. Each estimate is the linear slope of normalized sorting across within-year parental income deciles, where sorting between each dimension and occupational aspirations is normalized by the dimension’s entropy. Horizontal bars represent 95% percentile confidence intervals from 500 bootstrap replications.

value of anticipated earnings in a way that disproportionately affects sorting on ideology. But this is no longer an alternative to the arguments developed here. A diminishing-utility mechanism that operates markedly through expressive congruence is observationally and substantively equivalent to the claim that congruence carries greater weight under conditions of material security.

Log-Multiplicative Analyses for the Income Gradient

The mutual information analyses establish that the occupation–ideology association is stronger at higher parental incomes, but they cannot, by themselves, distinguish between two readings of that difference: higher-income students may exhibit a stronger version of the *same* occupation–ideology association, or they may hold a *different* association altogether, with different occupations carrying different political valences across income groups. This matters for interpretation: the luxury goods account implies that students across the income distribution share a common map of occupational reputations, not that they hold distinct political maps of the broad occupational structure.

To adjudicate between these readings, I turn to the log-multiplicative layer-effect model (Xie 1992). This model is useful since it separates two dimensions of the occupation–ideology association that mutual information combines: its structure across occupational categories and its overall strength

Table 2: Fit Statistics for the Log-Multiplicative Analyses

Model	Specification	G^2	df	% G^2 reduction	Δ
Independence	$\lambda + \lambda_o + \lambda_p + \lambda_g + \lambda_{og} + \lambda_{pg}$	82,108.0	78	0.0	7.7
Common Association	$M_0 + \psi_{op}$	8419.4	39	89.7	2.8
UNIDIFF	$M_0 + \phi_g \psi_{op}$	5253.7	38	93.6	1.9
Saturated	$M_0 + \psi_{opg}$	0.0	0	100.0	0.0

Notes: The table presents fit statistics for the four nested models. M_0 denotes the independence model. λ terms denote the grand mean, the main effects for occupation (o), political ideology (p), income group (g), and the occupation–group and ideology–group interactions. ψ_{op} denotes the occupation–ideology association and ϕ_g denotes the income-specific multiplier in the log-multiplicative (UNIDIFF) model. G^2 is model deviance; Δ , reported as percentage, is the share of cases that would have to be reallocated across cells for the fitted counts to match the observed counts.

within an income group. If individuals from different socioeconomic backgrounds act on a shared reputational map, the occupation-specific pattern should remain proportional across groups, with high-income students having a stronger version of the same association. If their reputational maps instead differ, allowing only the strength of a common structure to vary should fit poorly, because particular occupations would carry different political meanings across income groups. The model comparison therefore provides a direct test of whether the income gradient reflects an amplification of a shared association or, instead, a difference in the association’s occupational structure.

To evaluate these competing alternatives, I focus on liberals and conservatives only: I cross-tabulate occupational aspirations and political ideology within the bottom and top quartiles of within-year household income and I fit four nested models to these table counts:

1. **M0:** an independence model, in which occupations and ideology are unrelated within income;
2. **M1:** a common association model, which adds a single occupation–ideology association shared by both income groups;
3. **M2:** a log-multiplicative model (UNIDIFF), which retains the common association but allows its overall strength to differ by income group;
4. **M3:** a saturated model, in which each income group has its own occupation–ideology map.

Table 2 provides the fit statistics. I find that moving from the independence model to the common-association model reduces the deviance by 89.7% and lowers the dissimilarity index from 7.7% to 2.8%, suggesting that the occupation–ideology association is organized around a broadly common structure across the two income groups. Allowing the strength of that association to vary through the UNIDIFF specification improves the fit further, increasing the deviance reduction to 93.6% and lowering Δ to 1.9%. While this is substantively strong, the analyses provide only qualified support for Hypothesis 4: relative to Model 1, the log-multiplicative model explains 38% of the remaining




Figure 9: Occupation-Level Political Leans Across Income Quartiles

Notes: The figure presents, for each occupational category, the distribution of ideological sorting in occupational aspirations, computed as conservative share minus the liberal share, among students in the bottom quartile (horizontal axis) and the top quartile (vertical axis), pooled across 1976–2010. Each axis is centered at the corresponding quartile’s overall conservative–liberal balance, so values express an occupation’s lean relative to the average student in that quartile. Point size is proportional to occupation size. The dashed line indicates equal leans in the two quartiles; the solid line is a size-weighted linear fit through the origin. Participants who are “undecided” in their career aspirations are excluded.

deviance, compared to the saturated baseline. Given high sample size, we would, of course, expect that even small departures produce large deviance statistics, though this does not change the fact that Model 2 leaves more than half of the remaining deviance from Model 1 unexplained.

I investigate these mixed findings by analyzing the occupational sorting across the top and bottom quartiles. Figure 9 presents the results from this analysis, which contextualizes the results in Table 2. Consistent with Hypothesis 4, I find that the two income groups share a strongly similar political ordering of occupations. Inconsistent with Hypothesis 4, however, this pattern is not uniform, since the relative political positioning of several careers, most notably law and clergy, changes by income. These deviations slightly complicate a strictly uniform luxury-goods reading, though they are not strong enough to overturn it²⁶: the main ordering of occupations remains intact, and the departures are concentrated in a small number of careers. Taken together, the results provide qualified support for Hypothesis 4: high-income students appear to sort more strongly along a shared political map, even as the political composition of particular occupations varies somewhat by income.

²⁶The size-weighted fitted slope is $b = 1.30$, 95% CI = [1.02, 1.58], indicating that occupation-level political leans are, on average, approximately 30% stronger in the top income quartile as in the bottom quartile.

Discussion

Taken together, the results presented in this article support the idea that ideology is a consequential force in the formation of occupational aspirations, though its effect is unequally distributed across the stratification system. Using survey data on college freshmen entering U.S. institutions between 1976 and 2010, I showed that political ideology systematically structures occupational aspirations before labor market entry. I further showed that this association is not reducible to prior observable characteristics. Most consequentially, I showed that political sorting is significantly stronger among individuals from high-income families, proposing that the latitude to act on one's political ideology is itself stratified, making personal ideology a luxury good for occupational aspirations.

These findings suggest that political worldviews shape which careers individuals aspire to pursue well before workplace socialization or employer-level sorting, and that selection is a consequential mechanism of political stratification in economic life, complementing existing work that documents ideological sorting in the labor market setting (e.g., [Chinoy and Koenen 2024](#)). However, the results also suggest that this self-selection process is not uniform. The fact that ideology matters more for higher-income students means that the weight ideology carries in aspiration formation depends on an actor's position in the stratification system. Importantly, this implies that high-income students may structure the political character of occupations more so than low-income students.

Theoretical and Empirical Limitations

One substantive limitation of this article is the causal status of political ideology in the formation of occupational preferences. Since ideology and aspirations are measured simultaneously at college entry, I cannot rule out the possibility that students form preferences first and then adopt politically congruent positions. Such *anticipatory socialization*, the notion that committing to a career leads one to align their political commitments with their perceptions of what the proper political position is for that career, is a plausible alternative, and one path that this article cannot adjudicate. It is thus possible that students, anticipating cultural matching in hiring ([Rivera 2012](#)) or deciding that their political self-understandings should capture what occupations signal ([Akerlof and Kranton 2000](#)), shift their political views accordingly. That said, this scenario does not straightforwardly account for the income gradient, since the tendency to align one's politics with one's intended career would itself have to be stronger among higher-income students. This is not implausible: if higher-income students are more invested in their occupations, they might also be more apt to adjust their politics toward them. Hence, this study cannot adjudicate the *dynamics* of political alignment.

A related limitation is the perceptual basis of ideological congruence. The luxury-good mechanism assumes that students broadly share a common map of ideological reputations and differ primarily

in the weight they place on matching those reputations. The analyses have only qualified support for this assumption: the general ordering of occupations is similar across income, but the relative positioning of several careers varies with family income, suggesting that occupational reputations may not be perceived identically across the stratification system. Even if the reputations are shared, moreover, higher-income students may perceive its political orientations more clearly or with more confidence. Such differential perceptual clarity would produce stronger occupation–ideology sorting without requiring that political congruence carry greater weight in occupational choice. Since the survey does not directly measure students’ perceptions of occupational reputations, this study cannot fully resolve whether the income gradient arises from stronger preferences for congruence, sharper perceptions of occupations or some combination of these two mechanisms.

Several additional limitations concern the scope of the evidence. First, the central analysis captures occupational aspirations at college entry rather than the occupations students ultimately enter, and these aspirations may be revised as individuals encounter academic demands, financial constraints, hiring processes, or changing labor-market opportunities. Second, parental household income may be an imperfect measure of material security, since it does not directly capture family wealth, debt, expected parental support, or students’ confidence in the persistence of those resources. Moreover, the study is limited to students entering four-year colleges. Occupations individuals pursue largely without a college degree therefore fall outside this research, and the relationship between material security and ideological sorting may operate differently among non-college youth.

Implications for Theory and Research

Despite these limitations, these findings have several theoretical implications. I show that a meaningful share of the political composition of occupations is established before individuals enter the workforce. Existing accounts have documented the extent and magnitude of political sorting across firms and occupations ([Chinoy and Koenen 2024](#)), though leaving the question of how early in the life course these processes begin unanswered (but see [Meriläinen and Mitrunen 2025](#)). This article argues that firm-level processes operate on a population that is already ideologically sorted along occupational lines. While the evidence presented here does not exclude demand-side mechanisms, it introduces the development of occupational aspirations as a clear supply-side mechanism.

More broadly, these findings indicate that political ideology is a cultural force in the stratification of occupational life, functioning as an organizing principle through which people come to envision which careers are desirable in the first place ([Doepke and Zilibotti 2008](#)). In this sense, I argue that ideology is a structuring principle for “what people want” (see [Vaisey 2010](#)). At the same time, the luxury goods finding complicates any straightforward account of ideology as a universal force in preference formation. If political congruence is indeed an expressive good whose salience increases

with material security, it means that the politics of occupational choice relates to broader arguments about how expressive individualism is distributed unequally across the economic structure.

Ultimately, this article suggests that political ideology is not merely a correlate of where individuals end up in the labor market, but a consequential force that shapes economic outcomes. The tension I emphasized throughout this article, however, is that the capacity to act on such commitments is itself stratified. The consequence is a labor market in which the political *character* of occupations is disproportionately defined by those at the top of the income distribution. Whether such a dynamic deepens political divisions or dampens them is an unresolved question, but its implications for how we theorize the relationship between culture and inequality are, I believe, quite substantial.

References

- Acton, Riley, Emily Cook, and Paolo Ugalde Araya. 2025. “**Political Views and College Choices in a Polarized America**’.” SSRN, IZA Discussion Paper No. 18099.
- Akerlof, George A., and Rachel E. Kranton. 2000. “Economics and Identity.” *Quarterly Journal of Economics* 115(3):715–53. doi:[10.1162/003355300554881](https://doi.org/10.1162/003355300554881).
- Alonso-Villar, Olga, Coral Del Rio, and Carlos Gradin. 2012. “The Extent of Occupational Segregation in the United States: Differences by Race, Ethnicity, and Gender.” *Industrial Relations: A Journal of Economy and Society* 51(2):179–212. doi:[10.1111/j.1468-232X.2012.00674.x](https://doi.org/10.1111/j.1468-232X.2012.00674.x).
- Arold, Benjamin W. 2024. “Evolution Vs. Creationism in the Classroom: The Lasting Effects of Science Education.” *The Quarterly Journal of Economics* 139(4):2331–75. doi:[10.1093/qje/qjae019](https://doi.org/10.1093/qje/qjae019).
- Åslund, Olof, and Oskar Nordström Skans. 2009. “How to Measure Segregation Conditional on the Distribution of Covariates.” *Journal of Population Economics* 22(4):971–81. doi:[10.1007/s00148-008-0189-4](https://doi.org/10.1007/s00148-008-0189-4).
- Bailey, Martha, and Susan Dynarski. 2011. “**Gains and Gaps: Changing Inequality in U.S. College Entry and Completion**’.” Cambridge, MA: National Bureau of Economic Research.
- Becker, Gary S. [1964] 1994. *Human Capital: A Theoretical and Empirical Analysis, with Special Reference to Education*. Chicago: University of Chicago Press.
- Bermiss, Y. Sekou, and Rory McDonald. 2018. “Ideological Misfit? Political Affiliation and Employee Departure in the Private-Equity Industry.” *Academy of Management Journal* 61(6):2182–2209. doi:[10.5465/amj.2016.0817](https://doi.org/10.5465/amj.2016.0817).
- Blau, Peter M., and Otis Dudley Duncan. 1967. *The American Occupational Structure*. New York: Free Press.
- Brown, Jacob R., and Ryan D. Enos. 2021. “The Measurement of Partisan Sorting for 180 Million Voters.” *Nature Human Behaviour* 5(8):998–1008. doi:[10.1038/s41562-021-01066-z](https://doi.org/10.1038/s41562-021-01066-z).

- Bruch, Elizabeth E., and Robert D. Mare. 2012. "Methodological Issues in the Analysis of Residential Preferences, Residential Mobility, and Neighborhood Change." *Sociological Methodology* 42(1):103–54. doi:[10.1177/0081175012444105](https://doi.org/10.1177/0081175012444105).
- Chetty, Raj, Nathaniel Hendren, Patrick Kline, and Emmanuel Saez. 2014. "Where Is the Land of Opportunity? The Geography of Intergenerational Mobility in the United States." *The Quarterly Journal of Economics* 129(4):1553–623. doi:[10.1093/qje/qju022](https://doi.org/10.1093/qje/qju022).
- Chinoy, Sahil, and Martin Koenen. 2024. "'Political Sorting in the U.S. Labor Market: Evidence and Explanations'." Harvard University, Working Paper.
- Colonnelli, Emanuele, Valdemar Pinho Neto, and Edoardo Teso. 2025. "Politics at Work." *American Economic Review* 115(10):3367–3414. doi:[10.1257/aer.20240151](https://doi.org/10.1257/aer.20240151).
- DellaPosta, Daniel, Yongren Shi, and Michael Macy. 2015. "Why Do Liberals Drink Lattes?" *American Journal of Sociology* 120(5):1473–1511. doi:[10.1086/681254](https://doi.org/10.1086/681254).
- Doepke, Matthias, and Fabrizio Zilibotti. 2008. "Occupational Choice and the Spirit of Capitalism." *Quarterly Journal of Economics* 123(2):747–93. doi:[10.1162/qjec.2008.123.2.747](https://doi.org/10.1162/qjec.2008.123.2.747).
- Enke, Benjamin, Mattias K. Polborn, and Alex A. Wu. 2025. "Values as Luxury Goods and Political Behavior." *Journal of the European Economic Association* 23(5):1635–68. doi:[10.1093/jeea/jvaf014](https://doi.org/10.1093/jeea/jvaf014).
- Featherman, David L., and Robert Mason Hauser. 1978. *Opportunity and Change*. New York: Academic Press.
- Frake, Justin, Reuben Hurst, and Max Kagan. 2023. "'Political Segregation in the US Workplace'." SSRN, Working Paper.
- Goldstein, Yoav, and Matan Kolerman. 2025. "'Changing Minds: How Academic Fields Shape Political Attitudes'." SSRN, Working Paper.
- Gross, Neil. 2013. *Why Are Professors Liberal and Why Do Conservatives Care?* Cambridge: Harvard University Press.
- Haller, Archibald O., and Alejandro Portes. 1973. "Status Attainment Processes." *Sociology of Education* 46(1):51–91. doi:[10.2307/2112205](https://doi.org/10.2307/2112205).
- Hellerstein, Judith K., and David Neumark. 2008. "Workplace Segregation in the United States: Race, Ethnicity, and Skill." *The Review of Economics and Statistics* 90(3):459–77. doi:[10.1162/rest.90.3.459](https://doi.org/10.1162/rest.90.3.459).
- Hinich, Melvin J., and Michael C. Munger. 1998. *Ideology and the Theory of Political Choice*. Michigan: University of Michigan Press.
- Holland, John L. 1959. "A Theory of Vocational Choice." *Journal of Counseling Psychology* 6(1):35–45. doi:[10.1037/h0040767](https://doi.org/10.1037/h0040767).
- Hout, Michael. 1984. "Status, Autonomy, and Training in Occupational Mobility." *American Journal of Sociology* 89(6):1379–409. doi:[10.1086/228020](https://doi.org/10.1086/228020).
- Hout, Michael. 2018. "Americans' Occupational Status Reflects the Status of Both of Their Parents." *Proceedings of the National Academy of Sciences* 115(38):9527–32. doi:[10.1073/pnas.1802508115](https://doi.org/10.1073/pnas.1802508115).

- Inglehart, Ronald. 1997. *Modernization and Postmodernization: Cultural, Economic, and Political Change in 43 Societies*. Princeton: Princeton University Press.
- Kagan, Max, Justin Frake, and Reuben Hurst. 2025. “‘VRscores: A New Measure and Dataset of Workforce Politics Using Voter Registrations.’” SSRN, Working Paper.
- Kalleberg, Arne L. 1977. “Work Values and Job Rewards: A Theory of Job Satisfaction.” *American Sociological Review* 42(1):124–43. doi:10.2307/2117735.
- Kerckhoff, Alan C. 1976. “The Status Attainment Process: Socialization or Allocation?” *Social Forces* 55(2):368–81. doi:10.2307/2576228.
- Kitschelt, Herbert, and Philipp Rehm. 2014. “Occupations as a Site of Political Preference Formation.” *Comparative Political Studies* 47(12):1670–1706. doi:10.1177/0010414013516066.
- Lin, Nan, Walter M. Ensel, and John S. Vaughn. 1981. “Social Resources and Strength of Ties: Structural Factors in Occupational Status Attainment.” *American Sociological Review* 46(4):393–405. doi:10.2307/2095260.
- Lynn, Freda B., and George Ellerbach. 2017. “A Position with a View: Educational Status and the Construction of the Occupational Hierarchy.” *American Sociological Review* 82(1):32–58. doi:10.1177/0003122416671743.
- Lynn, Freda B., Yongren Shi, and Kevin Kiley. 2025. “Intersectional Group Agreement on the Occupational Order.” *Social Psychology Quarterly* 88(1):135–48. doi:10.1177/01902725241256378.
- Mas, Alexandre, and Amanda Pallais. 2017. “Valuing Alternative Work Arrangements.” *American Economic Review* 107(12):3722–59. doi:10.1257/aer.20161500.
- McFadden, Daniel. 1974. “Conditional Logit Analysis of Qualitative Choice Behavior.” Pp. 105–42 in *Frontiers in Econometrics*, edited by P. Zarembka. New York: Academic Press.
- Meriläinen, Jaakko, and Matti Mitrunen. 2025. “‘Communism in the Classroom: Long-Run Effects of an Experiment.’” PhilPapers, Working Paper.
- Mora, Ricardo, and Javier Ruiz-Castillo. 2011. “Entropy-Based Segregation Indices.” *Sociological Methodology* 41(1):159–94. doi:10.1111/j.1467-9531.2011.01237.x.
- Reardon, Sean F., and Glenn Firebaugh. 2002. “Measures of Multigroup Segregation.” *Sociological Methodology* 32(1):33–67. doi:10.1111/1467-9531.00110.
- Reny, Tyler T., Marcel F. Roman, Benjamin J. Newman, and David O. Sears. 2025. “Late Adolescents Entering College Intending a Career as Police Officers Hold More Right-Leaning Views Than Their Peers.” *Proceedings of the National Academy of Sciences* 122(32):e2500220122. doi:10.1073/pnas.2500220122.
- Reskin, Barbara F., Debra B. McBrier, and Julie A. Kmec. 1999. “The Determinants and Consequences of Workplace Sex and Race Composition.” *Annual Review of Sociology* 25:335–61. doi:10.1146/annurev.soc.25.1.335.
- Rivera, Lauren A. 2012. “Hiring as Cultural Matching: The Case of Elite Professional Service Firms.” *American Sociological Review* 77(6):999–1022. doi:10.1177/0003122412463213.

- Rosen, Sherwin. 1986. "[The Theory of Equalizing Differences](#)." Pp. 641–92 in *Handbook of Labor Economics*. Vol. 1, edited by O. C. Ashenfelter and R. Layard. Amsterdam: North-Holland.
- Schuessler, Alexander A. 2000. *A Logic of Expressive Choice*. Princeton: Princeton University Press.
- Scott, Ralph. 2022. "Does University Make You More Liberal? Estimating the Within-Individual Effects of Higher Education on Political Values." *Electoral Studies* 77:102471. doi:[10.1016/j.electstud.2022.102471](#).
- Selenko, Eva, Miriam Schilbach, Steven A. Brieger, Anahí Van Hootegem, and Hans De Witte. 2025. "The Political Consequences of Work: An Integrative Review." *Journal of Management* 51(6):2355–88. doi:[10.1177/01492063241301337](#).
- Sewell, William H., Archibald O. Haller, and Alejandro Portes. 1969. "The Educational and Early Occupational Attainment Process." *American Sociological Review* 34(1):82–92. doi:[10.2307/2092789](#).
- Spilerman, Seymour. 1977. "Careers, Labor Market Structure, and Socioeconomic Achievement." *American Journal of Sociology* 83(3):551–93. doi:[10.1086/226595](#).
- Vaisey, Stephen. 2010. "What People Want: Rethinking Poverty, Culture, and Educational Attainment." *The ANNALS of the American Academy of Political and Social Science* 629(1):75–101. doi:[10.1177/0002716209357146](#).
- Weeden, Kim A. 2002. "Why Do Some Occupations Pay More Than Others? Social Closure and Earnings Inequality in the United States." *American Journal of Sociology* 108(1):55–101. doi:[10.1086/344121](#).
- Weeden, Kim A., and David B. Grusky. 2005. "The Case for a New Class Map." *American Journal of Sociology* 111(1):141–212. doi:[10.1086/428815](#).
- Xie, Yu. 1992. "The Log-Multiplicative Layer Effect Model for Comparing Mobility Tables." *American Sociological Review* 57(3):380–95. doi:[10.2307/2096242](#).
- York, Hunter, Xi Song, and Yu Xie. 2025. "Gradationalism Revisited: Intergenerational Occupational Mobility Along Axes of Occupational Characteristics." *American Journal of Sociology* 130(4):976–1027. doi:[10.1086/733122](#).

Supplemental Materials for Ideology as a Luxury Good in Occupational Choice

Turgut Keskintürk

Table of Contents

Supplemental Materials A: Theoretical Model	S2
Supplemental Materials B: Data Sources and the Characteristics of the Sample	S4
Supplemental Materials C: Additional Analyses for Ideological Sorting	S9
C.1. Robustness to Marginals	S9
C.2. Robustness to Weighting	S13
C.3. The Distribution of Ideological Sorting Over Time	S15
C.4. Aspirations versus Realized Partisanship	S16
C.5. Analyses Using Political Beliefs and Preferences	S18
Supplemental Materials D: Additional Analyses for Ideology as a Luxury Good	S23
D.1. Robustness to Marginals	S23
D.2. Robustness to Weighting	S28
D.3. Robustness to Income Measurement	S29
D.4. Reproducing the Income Gradient Across the Study Window	S31
D.5. Students with Undecided Career Aspirations	S32
D.6. Robustness to the Composition of the Occupation Set	S34
D.7. Income Gradient Across Parental Education	S35
D.8. Robustness to Occupational Reproduction Across Generations	S37
D.9. Analyses Using Political Beliefs and Preferences	S38
References	S39

Supplemental Materials A: Theoretical Model

I present a simple utility model that formalizes “political congruence” as an occupational amenity (Rosen 1986), using a canonical ideological distance term widely used in spatial models of political choice (Hinich and Munger 1998). Let actor i ’s expected utility from occupation j simply be:

$$U_{ij} = [1 - \lambda(y_i)]E_{ij} + \lambda(y_i)C_{ij} + \epsilon_{ij} \quad (\text{A1})$$

Following Enke, Polborn, and Wu (2025), Equation A1 says that actor i ’s utility from occupation j is a weighted average of material (E) and non-material (C) rewards. Suppose that each actor i has a type (y_i, p_i) , with y_i denoting parental income and $p_i \in \mathbb{R}$ denoting i ’s ideological position.

We can specify the material component of this utility term E using expected earnings w_j as well as a vector of amenities a_j (e.g., working conditions, security, or commuting):

$$E_{ij} = w_j + \beta'_i a_j \quad (\text{A2})$$

Consistent with classic research on compensating differentials (Mas and Pallais 2017; Rosen 1986), the vector β_i captures actor-specific weights on such amenities. In addition to the material returns, however, occupations might also offer expressive returns through their political character. Suppose that each occupation has a perceived ideological reputation $\pi_j \in \mathbb{R}$. Then:

$$C_{ij} = -\frac{1}{2}\gamma(p_i - \pi_j)^2 \quad (\text{A3})$$

where “congruence” is defined via symmetric quadratic loss as in spatial models of political choice and γ is a scaling parameter converting ideological distance into commensurate material units.

Connecting material and non-material returns, there is a weight function $\lambda(\cdot) : \mathbb{R} \rightarrow [0, 1]$, which is strictly increasing in parent income y . As actors become materially more secure, they place greater weight on non-material considerations, including congruence, relative to their material returns.

If we substitute E and C to Equation A1, we get the final utility form:

$$U_{ij} = [1 - \lambda(y_i)](w_j + \beta'_i a_j) - \frac{1}{2}\lambda(y_i)\gamma(p_i - \pi_j)^2 + \epsilon_{ij} \quad (\text{A4})$$

It is worth being explicit about what $\lambda(y)$ represents. In Equation A4, I model the income gradient as *taste heterogeneity*: as material security rises, actors place genuinely greater weight on expressive, identity-relevant returns relative to material ones (Inglehart 1997). As argued in the main text, an

observationally equivalent gradient may arise from the utility term itself: if the marginal utility of material returns is diminishing, the *relative* weight on congruence rises with income even when the underlying taste for congruence is fixed. I follow the taste-based formulation here because it most directly expresses this paper's claim that ideological congruence is an *expressive good*.

Let the discrete choice variable J denote the actor's chosen occupation, where an individual selects $j^* = \arg \max_j U_{ij}$. Under this decision rule, the probability that actor i selects occupation j rises in U_{ij} , so the responsiveness of i 's choice to an occupation's political reputation π_j is governed by the weight $\lambda(y_i)$. As shown in Equation A5 below, the effect of an occupation's congruence advantage on relative utility, $\Delta C_i = C_{ij} - C_{ij'}$, increases as a function of their parental income y_i :

$$\frac{\partial(U_{ij} - U_{ij'})}{\partial \Delta C_i} = \lambda(y_i), \quad \frac{\partial}{\partial y_i} \left[\frac{\partial(U_{ij} - U_{ij'})}{\partial \Delta C_i} \right] = \lambda'(y_i) > 0 \quad (\text{A5})$$

Aggregated across the population, the higher-income actors of a given ideology concentrate their choices more tightly on occupations whose reputation π_j matches their position p_i . Therefore, the occupational distribution conditional on ideology, $\Pr(J | p)$, departs further from the marginal $\Pr(J)$ under this model. The mutual information M is precisely the measure for this departure:

$$\frac{\partial M(J; P | y)}{\partial y} > 0 \quad (\text{A6})$$

In words, the mutual information M between people's political ideology P and chosen occupation J is increasing in income, which is what the luxury goods hypothesis explicitly proposes.

Supplemental Materials B: Data Sources and the Characteristics of the Sample

The Freshman Survey (TFS) is a large-scale survey study, designed to capture first-year students at U.S. institutions of higher education granting baccalaureate-level degrees or above, as described in the *Integrated Postsecondary Education Data System* at the Department of Education. Starting with its inception in the 1960s, more than 2,000 organizations have participated in these surveys at different time points, with more than 200,000 students in each year. The breadth, granularity, and timing of this data make TFS an ideal candidate for an empirical examination of preference formation.

I made several decisions while processing these surveys for analyses. First, I excluded participants who were aged 21 or older at the time of the survey ($\approx 5.7\%$ of the sample) to focus on young adults. Second, I excluded surveys completed before 1976 because the survey item used to elicit students' occupational preferences had substantial revisions during this window. Third, I only included data with complete responses to the occupation and ideology items (leading to 12% decrease). Table B1 presents the basic descriptive statistics for weighted observations, unweighted observations, and observations with missing weights, showing no meaningful differences across these groups.

Table B1: Descriptive Statistics in TFS

	Unweighted	Weighted	Weight Missing
Characteristic	N = 9,506,408	N = 6,696,391	N = 2,810,017
Gender			
Male	45.3%	46.7%	45.0%
Female	54.7%	53.3%	55.0%
Age	18.27 (0.55)	18.23 (0.50)	18.39 (0.66)
Race			
White	79.0%	78.7%	77.3%
Black	6.9%	8.4%	8.1%
Hispanic	3.4%	3.3%	4.1%
Asian	5.3%	4.4%	4.9%
Other	5.4%	5.2%	5.7%
Religion			
Protestant	42.4%	44.8%	43.4%
Catholic	32.0%	31.2%	31.1%
Jewish	3.8%	3.0%	3.3%
Other	6.8%	7.1%	7.8%
None	14.9%	14.0%	14.3%
Political Ideology			
Liberal	28.6%	27.1%	27.3%
Moderate	49.0%	50.9%	50.9%
Conservative	22.4%	22.0%	21.7%

(continued)

	Unweighted	Weighted	Weight Missing
Characteristic	N = 9,506,408	N = 6,696,391	N = 2,810,017
Type of Institution			
2-Year College	2.5%	0.0%	8.4%
4-Year College	52.0%	55.4%	51.0%
University	45.5%	44.6%	40.6%
Historically Black College	2.6%	4.1%	3.5%

TFS uses 26 stratification units to group the participating organizations based on institutional type (e.g., four-year colleges or universities), organizational race, selectivity, and financial control. The participation, however, is voluntary at the institution level, so these surveys are not nationally representative. To address this limitation, TFS provides sampling weights that reflect the stratification design and applies the requirement that the institutions survey at least 65% of their students.¹

The 2016 *Cooperative Institutional Research Program* (CIRP) report presents long-standing trends in TFS, as well as more details about the research methodology, the construction of sampling weights, and comparability across time (see [Eagan et al. 2016](#)). The original survey instruments, codebooks, and institution-level participation status across years are accessible from the electronic [Instruments and Codebooks](#) page at the HERI website. The original variable names, their corresponding data files, as well as the harmonization schemes I employed can be found in the reproduction repo.

As discussed in the manuscript, I use a harmonized list of occupational intentions for the analyses, and I demonstrate the distribution of these aspirations, with and without weights, in Table B2.

Table B2: Occupational Aspirations in the TFS

	Unweighted	Weighted	Weight Missing
Characteristic	N = 9,506,408	N = 6,696,391	N = 2,810,017
Occupational Preferences			
Accountant	3.6%	4.0%	4.0%
Actor or Entertainer	1.4%	1.3%	1.4%
Administrative Assistant	0.7%	0.6%	1.1%
Artist	1.9%	1.9%	2.3%
Business Manager/Executive	9.0%	9.0%	8.9%
Business Owner/Entrepreneur	2.7%	2.8%	2.9%
Clergy	0.4%	0.4%	0.5%
Clinical Psychologist	1.6%	1.6%	1.6%
College Faculty	0.5%	0.4%	0.5%
Computer Programmer/Developer	2.4%	2.7%	2.9%

¹This requirement is the main reason why there is high missingness in sampling weights provided by TFS.

(continued)

Characteristic	Unweighted	Weighted	Weight Missing
	N = 9,506,408	N = 6,696,391	N = 2,810,017
Dentist/Orthodontist	0.8%	0.9%	0.8%
Dietitian/Nutritionist	0.4%	0.4%	0.4%
Elementary School Teacher	4.1%	4.8%	4.5%
Engineer	7.2%	7.3%	7.3%
Farmer or Forester	1.4%	1.4%	1.4%
Foreign Service Worker	1.0%	0.8%	0.8%
Homemaker	0.2%	0.2%	0.2%
Interior Designer	0.4%	0.4%	0.5%
Lawyer/Judge	5.4%	4.8%	4.5%
Medical Doctor/Surgeon	7.0%	6.1%	5.3%
Medical/Dental Assistant	0.8%	0.8%	0.9%
Military	1.1%	1.2%	1.1%
Musician	1.4%	1.4%	1.4%
Natural Resource Specialist	1.1%	1.0%	1.2%
Optometrist	0.3%	0.3%	0.2%
Other	7.8%	7.9%	8.6%
Pharmacist	1.2%	1.3%	1.1%
Protective Services	0.9%	1.1%	1.2%
Registered Nurse	2.8%	3.1%	3.5%
Research Scientist	2.1%	2.0%	1.9%
Retail Sales	1.0%	1.1%	1.1%
School Counselor	0.4%	0.4%	0.4%
School Principal/Superintendent	0.1%	0.1%	0.1%
Secondary School Teacher	3.4%	3.9%	3.4%
Skilled Trades	0.5%	0.4%	0.9%
Social Worker	1.4%	1.4%	1.5%
Therapist	2.6%	2.9%	2.6%
Undecided	13.6%	13.2%	12.2%
Urban Planner/Architect	1.2%	1.1%	1.4%
Veterinarian	1.2%	1.2%	1.2%
Writer/Producer/Director	2.8%	2.6%	2.6%

The distributions of occupational aspirations and ideology across the study window are presented in Figure B1 and Figure B2, respectively, highlighting several trends among college students.

Undecided	11.3	11.7	13.3	12.2	13.3	16.5	15.8
Business Manager/Executive	8.9	11.3	11.8	7.5	8.7	7.9	7.8
Other	6.1	5.7	6.7	6.5	8.7	9.8	9.6
Engineer	9.2	11.4	9.0	8.2	0.2	7.0	7.4
Medical Doctor/Surgeon	5.9	6.3	5.8	8.6	7.7	6.6	7.6
Lawyer/Judge	6.3	6.0	6.8	5.9	4.4	4.7	4.2
Elementary School Teacher	3.5	2.8	4.0	4.4	5.5	4.7	3.6
Accountant	5.4	5.3	4.8	3.6	2.5	2.1	2.3
Secondary School Teacher	2.6	1.9	3.1	3.9	4.0	4.1	3.8
Writer/Producer/Director	2.6	2.7	3.1	2.5	2.9	2.9	3.1
Registered Nurse	3.6	3.1	1.6	2.7	2.0	2.9	4.0
Business Owner/Entrepreneur	2.0	2.5	3.3	2.3	2.7	3.1	3.0
Therapist	2.2	2.1	2.2	2.9	3.2	2.4	2.9
Computer Programmer/Developer	2.9	5.6	2.2	2.0	0.7	2.7	1.6
Research Scientist	2.9	2.2	2.2	2.3	1.1	2.0	2.3
Artist	1.8	1.6	1.8	1.7	1.9	2.2	2.1
Clinical Psychologist	1.3	1.3	1.8	2.0	1.8	1.6	1.4
Farmer or Forester	0.7	0.5	0.3	0.3	7.1	0.2	0.2
Musician	1.5	1.2	1.2	1.3	1.5	1.5	1.6
Social Worker	2.1	1.1	1.0	1.4	2.3	0.9	0.9
Actor or Entertainer	1.2	1.0	1.2	1.3	1.7	1.7	1.3
Pharmacist	0.8	0.7	0.9	1.2	1.0	1.7	2.0
Urban Planner/Architect	1.4	1.1	1.4	1.3	1.1	1.1	1.0
Veterinarian	1.3	1.0	0.8	1.7	1.2	1.1	1.2
Military	1.3	1.3	1.1	0.9	0.9	1.0	1.2
Natural Resource Specialist	1.1	0.5	0.4	0.6	4.0	0.2	0.3
Retail Sales	1.0	1.2	1.3	0.8	1.0	1.0	0.9
Foreign Service Worker	0.9	1.2	1.7	1.1	0.3	0.9	1.2
Protective Services	0.9	0.7	0.7	1.1	1.0	1.0	1.0
Dentist/Orthodontist	1.3	0.8	0.6	0.6	0.4	0.9	1.2
Medical/Dental Assistant	1.3	0.8	0.3	3.0	0.2	0.2	0.2
Administrative Assistant	1.2	0.8	0.5	0.5	0.6	0.6	0.6
Skilled Trades	0.7	0.6	0.4	0.9	0.4	0.3	0.3
College Faculty	0.4	0.4	0.6	0.8	0.0	0.6	0.6
School Counselor	0.3	0.2	0.3	0.5	1.1	0.3	0.3
Clergy	0.5	0.4	0.3	0.4	0.5	0.5	0.3
Interior Designer	0.6	0.5	0.6	0.3	0.1	0.5	0.4
Dietitian/Nutritionist	0.5	0.3	0.2	0.2	0.7	0.3	0.4
Optometrist	0.3	0.2	0.3	0.3	0.2	0.2	0.3
Homemaker	0.2	0.1	0.1	0.1	0.8	0.1	0.1
School Principal/Superintendent	0.0	0.0	0.0	0.1	0.3	0.1	0.0
	1976-1980	1981-1985	1986-1990	1991-1995	1996-2000	2001-2005	2006-2010

Figure B1: The Distribution of Occupational Aspirations Over Time

Notes: The figure presents the marginal distribution of occupational aspirations within five-year periods, with percents summing to 100% within each period, and darker cells indicating larger shares. Occupations are ordered by their pooled prevalence across the full study window. The “Other” and “Undecided” categories are included in the figure.




Figure B2: The Distribution of Political Ideology Over Time

Notes: The figure presents the share of students identifying as liberal, moderate, and conservative across survey years, as weighted and unweighted series.

Supplemental Materials C: Additional Analyses for Ideological Sorting

In this section, I present several specification checks for results presented in the article.

C.1. Robustness to Marginals

Check 1. The expected M under independence is effectively zero in every year, which suggests that the observed values are not due to finite-sample noise and are comparable with minimal bias.

Check 2. The “effective” number of occupational categories, which I calculated as the exponential of Shannon entropy, remained largely stable across this window, ranging from 22 to 28, while the entropy of occupational distributions (ranging from 3.1 to 3.3) as well as political ideology (ranging from 1 to 1.1) minimally varied over time. Taken together, these values suggest that the trajectories in political sorting are not just mechanical artifacts of shifts in the marginal distributions. To show an alternative specification, I decompose changes in M into three pathways in Figure C1², showing that the divergence between liberals and conservatives is a structural change, rather than an artifact of changing marginals. I also show that analyses with normalized M values—which normalize the observed M scores with the entropy of occupations (Figure C2) and the entropy of ideology groups and relevant benchmark categories (Figure C3)—provide the same substantive conclusions.³

²I use an informative visualization strategy introduced in Sobchuk and Beheim (2025) for representing these pathways.

³Using normalized M scores for benchmarking provides consistent findings for 2010. Figure C2 suggests that political sorting is 23% as large as gender sorting, and 79% as large as racial sorting. Figure C3 suggests that political sorting is 15% as large as gender sorting, and 82% as large as racial sorting.





Figure C1: The Decomposition of Temporal Trajectories in Ideological Sorting

Notes: The figure presents the decomposition of consecutive year-to-year changes in political sorting into (1) change due to occupation marginals, (2) change due to ideology marginals, and (3) change due to the occupation-ideology relation. For each adjacent pair of M , the change in M is decomposed with a Shapley procedure. Bars are represented as floating contributions (stacking positive and negative contributions), and the black line presents the observed M values. Panel A (top row): all three ideology categories and Panel B (bottom row): liberals and conservatives only.

Ideology



Benchmarks





Figure C2: Temporal Trajectories in Ideological Sorting, Normalized by Occupational Entropy

Notes: The figure presents changes in ideological sorting over time, and benchmarks ideological sorting against sorting by gender and sorting by race. Each panel represents yearly M values between occupational aspirations and the relevant grouping variable (described on the bottom-right corner of each panel), as weighted and unweighted series. Participants who are “undecided” in their career aspirations are excluded. The shaded band marks 1995–2000 to highlight temporary changes in the occupational preference instrument in the surveys. “Ideologues” in Panel C refer to an alternative coding where liberals and conservatives are combined into one group for comparisons with moderates. I normalize the mutual information values by the entropy of occupations within each year to calculate normalized mutual information.

Ideology



Benchmarks



Figure C3: Temporal Trajectories in Ideological Sorting, Normalized by Group Entropy

Notes: The figure presents changes in ideological sorting over time, and benchmarks ideological sorting against sorting by gender and sorting by race. Each panel represents yearly M values between occupational aspirations and the relevant grouping variable (described on the bottom-right corner of each panel), as weighted and unweighted series. Participants who are “undecided” in their career aspirations are excluded. The shaded band marks 1995–2000 to highlight temporary changes in the occupational preference instrument in the surveys. “Ideologues” in Panel C refer to an alternative coding where liberals and conservatives are combined into one group for comparisons with moderates. I normalize the mutual information values by the entropy of the group within each year to calculate normalized mutual information.

C.2. Robustness to Weighting

The decomposition analysis presented in the main article employed unweighted estimates. In Figure C4, I show the same analysis with weighted estimates, with substantively similar results. I also reproduce occupational contributions with weights in similar fashion, presented in Figure C5.




Figure C4: Weighted Decomposition of Ideological Sorting by Sociodemographic Characteristics

Notes: The figure presents a decomposition of ideological sorting in occupational preferences by sociodemographic background characteristics. Baseline refers to observed M, while other rows present residual M after sequentially expanding the set of predictors used to construct counterfactual ideology distributions. Panel A uses all three ideology categories, while Panel B restricts the sample to liberals and conservatives. Missing values are included in the estimations.




Figure C5: Weighted Occupation-Specific Contributions to Ideological Sorting

Notes: The figure presents the decomposition of ideological sorting into occupation-specific contributions. Each cell represents an occupation's share of the aggregate mutual information M in the corresponding period, with shares summing to 100% within each period, and darker cells indicating larger contributions. Participants who are "undecided" in their career aspirations, as well as those in the residual "Other" category, are excluded from the decomposition.

C.3. The Distribution of Ideological Sorting Over Time

Figure C6 traces yearly composition of ideology for each career across the study window.




Figure C6: Occupation-Level Ideological Distributions Over Time

Notes: The figure presents yearly composition values (conservative share minus liberal share) for each career. Each facet presents a single occupation, with facets ordered from most conservative (top-left) to most liberal (bottom-right). The shaded band marks 1995–2000 to highlight temporary changes in the occupational preference instrument.

C.4. Aspirations versus Realized Partisanship

A natural question about occupational aspirations is whether they have any systematic relationship to the political composition of the labor market itself. To investigate this, I use *VR-Scores*, a dataset that estimates the partisan composition of occupations by linking the L2 voter files to employment histories from the *Revelio Labs* (Kagan, Frake, and Hurst 2025). *VR-Scores* provide occupation-level estimates of two-party registration—the share of workers registered as Republican minus the share registered as Democrats—for detailed O*NET occupational titles, with coverage beginning in 2012. I use the 2012 panel, the closest available time window to the final year of the TFS series.

I proceed in three steps. First, I compute occupation-level ideological differentiation in aspirations using the 2010 TFS wave, excluding the “Undecided” and “Other” categories. Second, I match the occupational titles in *VR-Scores* to TFS career categories using a crosswalk, which matches 646 titles onto 36 TFS careers. For each career, I then aggregate partisan composition using an employment-weighted average and compute the 25th and 75th percentiles of these composition values to encode heterogeneity in the mapping process. Third, I calculate the correlation between aspiration-based differentiation and realized partisanship. I then bootstrap this entire pipeline (resampling O*NET titles within careers over 1,000 draws) to obtain confidence intervals for the correlation.

Figure C7 presents the findings from this exercise: Panel A presents each career’s realized partisan composition against its aspiration-based differentiation; Panel B presents the bootstrap distribution of the correlation; and Panel C displays each career’s employment-weighted margin alongside its interquartile range across constituent titles. I show an overall correlation of $\rho = 0.88$ (with 95% CI: $[0.80, 0.90]$): careers that disproportionately attract conservative students are, almost without exception, the careers disproportionately staffed by registered Republicans. This association does not establish that aspirations produce the realized composition, though it still shows that the political organization of aspirations and the political organization of work mirror one another closely.




Figure C7: Ideological Sorting in Aspirations and Realized Partisan Composition

Notes: The figure compares ideological differentiation in aspirations with the realized partisan composition of the corresponding occupations. Panel A: realized partisan margins against aspiration-based differentiation, with point size proportional to employment. Panel B: the bootstrap distribution of the Spearman correlation from 1,000 draws, resampling O*NET titles within each career to incorporate uncertainty in the matching. Panel C: employment-weighted margins by career, with segments marking the employment-weighted interquartile range across constituent O*NET titles.

C.5. Analyses Using Political Beliefs and Preferences

It is possible to argue that some of the time trends reflect better *sorting* into ideological labels, rather than stronger politicization of occupational aspirations. Considering that political ideology is now more constrained than before (e.g., Baldassarri and Gelman 2008; Levendusky 2009), the temporal trends I observe in the article might reflect a more general sorting trajectory. In Figure C8, I present alternative analyses that employ substantive policy positions to predict occupational sorting. These findings show that these issue positions are predictive of occupational aspirations, yet many of the effects are stable across time. Considering that the “increasing constraint” hypothesis is supported in this data as well (Figure C9), it is likely that a share of these trends reflects increased ideological constraint over the years rather than substantive politicization of occupational aspirations.

To explore issue-occupation coupling better, I analyzed average positions at the level of issues and occupations. Figure C10 demonstrates a heatmap of standardized vocation-level deviations from the issue mean on each issue where positive values (red) indicate more conservative positions and negative values (blue) more liberal ones. The figure presents strong descriptive evidence on how political differentiation in occupational aspirations spans multiple issue domains, with careers that skew conservative on one policy tending to do so on others as well. There are also several notable examples of group-specific ideological positioning: students aspiring to be a member of the clergy, for instance, are much more conservative on the legality of marijuana, homosexuality, and abortion, while being more liberal on capital punishment. Similarly, a distinctive issue position among those who would like to become a business owner/entrepreneur is raising taxes on the wealthy.⁴

⁴In Table C1, I show the response distributions and wordings for issue items used in the analyses.




Figure C8: Issue-Specific Ideological Sorting in Occupational Preferences

Notes: The figure presents how informative political issue positions are of occupational preferences over time, separately for each issue. The solid lines represent unweighted mutual information values, while dashed lines represent weighted values. All issue positions are measured with four response categories, including strongly disagree, disagree, agree and strongly agree. Complete item wording for each issue domain is presented in Table C1.




Figure C9: Ideology-Issue Alignment Over Time

Notes: The figure presents how strongly political ideology is associated with policy issue responses over time, separately for each issue domain. For each issue and year, alignment is measured as mutual information between ideology group and issue category, with solid lines representing unweighted mutual information values and dashed lines representing weighted values. All issue positions are measured with four response categories, including strongly disagree, disagree, agree and strongly agree. Complete item wording for each issue domain is presented in Table C1.




Figure C10: Issue Positions Across Occupations

Notes: The figure presents occupation-level standardized deviations on twelve issue items, computed from unweighted data pooled across all available survey years for each issue. Positive values (darker red) indicate that actors hold more conservative positions than the overall population, while negative values (darker blue) indicate more liberal positions.

Table C1: Response Distributions Across Political Issues

Survey Item	Strongly Disagree	Disagree	Agree	Strongly Agree	N
A national health care plan is needed to cover everybody's medical costs.	10.8	23.4	42.9	22.9	5,934,496
Abortion should be legal.	25.3	16.0	28.8	29.9	8,788,440
It is important to have laws prohibiting homosexual relationships.	37.4	30.1	18.1	14.4	8,787,850
Marijuana should be legalized.	40.8	25.6	22.6	10.9	9,054,474
Racial discrimination is no longer a major problem in America.	38.2	43.5	15.6	2.7	5,795,891
The activities of married women are best confined to the home and family.	53.7	25.1	14.6	6.6	7,585,153
The death penalty should be abolished.	34.7	35.6	17.3	12.4	8,284,648
The federal government is not doing enough to control environmental pollution.	2.4	15.4	48.5	33.7	6,746,425
The federal government should do more to control the sale of handguns.	6.2	13.5	38.2	42.1	5,983,517
The federal government should raise taxes to reduce the deficit.	23.8	47.0	24.8	4.3	4,323,934
There is too much concern in the courts for the rights of criminals.	7.8	28.4	43.4	20.4	8,343,882
Wealthy people should pay a larger share of taxes than they do now.	12.1	25.2	37.1	25.6	7,813,807

Supplemental Materials D: Additional Analyses for Ideology as a Luxury Good

In this section, I present several specification checks for results presented in the article.

D.1. Robustness to Marginals

Check 1. It might be that the gradient reflects differences in the entropy of occupational aspirations or ideological distributions across deciles. Substantively, the former would suggest that higher-SES students face broader occupational choice sets⁵, which mechanically gives more room for ideological differences to generate sorting, while the latter suggests group differences in ideology. I present evidence that normalizing M by within-decile occupation entropy (Figure D1) or ideology entropy (Figure D2) produces a comparable set of findings, ruling out these potential interpretations.

Check 2. In Figures D3 and D4, I show career sorting by gender and race, this time by normalizing them with the entropy of occupations (Figure D3) or the entropy of sex or race (Figure D4). Once again, the substantive conclusions in the main article hold under both of these specifications.

⁵Put differently, higher-SES students may aspire to a wider range of occupations. If the gradient survives normalization by within-decile occupational entropy, that would give strong evidence against this composition-based story.




Figure D1: Income Gradient in Ideological Sorting, Normalized by Occupational Entropy

Notes: The figure presents mutual information M between political ideology and occupational aspirations across household income, which I calculate as within-year income rank across 1976-2010. The estimates include bootstrap confidence intervals from 1,000 draws, as well as a linear fit to describe the income gradient. Panels A and B: all three ideology categories. Panels C and D: liberals and conservatives only. I normalize the mutual information values by the entropy of the occupations within each decile to calculate normalized mutual information.




Figure D2: Income Gradient in Ideological Sorting, Normalized by Ideology Entropy

Notes: The figure presents mutual information M between political ideology and occupational aspirations across household income, which I calculate as within-year income rank across 1976-2010. The estimates include bootstrap confidence intervals from 1,000 draws, as well as a linear fit to describe the income gradient. Panels A and B: all three ideology categories. Panels C and D: liberals and conservatives only. I normalize the mutual information values by the entropy of the group within each decile to calculate normalized mutual information.




Figure D3: Gender and Racial Sorting, Normalized by Occupational Entropy

Notes: The figure presents mutual information M between gender and occupational aspirations (top panels), and race and occupational aspirations (bottom panels), across household income, which I calculate as within-year income rank across 1976-2010. The columns present unweighted and weighted estimates. The estimates include bootstrap confidence intervals from 1,000 draws, as well as a quadratic fit to describe the income gradient. I normalize the mutual information values by the entropy of the occupations within each decile to calculate normalized mutual information.




Figure D4: Gender and Racial Sorting, Normalized by Group Entropy

Notes: The figure presents mutual information M between gender and occupational aspirations (top panels), and race and occupational aspirations (bottom panels), across household income, which I calculate as within-year income rank across 1976-2010. The columns present unweighted and weighted estimates. The estimates include bootstrap confidence intervals from 1,000 draws, as well as a quadratic fit to describe the income gradient. I normalize the mutual information values by the entropy of the group within each decile to calculate normalized mutual information.

D.2. Robustness to Weighting

The attribute decomposition analysis presented in the article used unweighted estimates. In Figure D5, I show the same analysis with weighted estimates, with substantively comparable findings.




Figure D5: Weighted Income Gradient in Ideological Sorting

Notes: The figure presents observed and residual mutual information between ideology and aspirations across parental income deciles in the 1989–1990 subsample. Residual M is computed using the same decomposition strategy described in the main paper, now applied within each income decile using the thirteen standardized career-attribute ratings. Panel A uses all three ideology categories, while Panel B restricts the sample to liberals and conservatives.

D.3. Robustness to Income Measurement

The analyses use within-year income decile ranks, which removes inflation and category-redesign artifacts, but obscures absolute purchasing power. Figure D6 shows that the gradient holds when income is instead measured as real parental income in 2010 dollars. Of course, a note on midpoint choices is warranted here. The only consequential researcher decision in constructing real income concerns the open-ended top categories (e.g., “\$150,000 or more”), to which I assign midpoints at 1.5 times the threshold, a common convention in income harmonization. Reasonable alternatives include using the thresholds themselves, doubling them, or assigning Pareto-implied conditional means estimated from the upper tail of the distribution. For the main analyses, these choices are immaterial by construction. For the absolute-income analyses presented in Figure D6, the alternatives matter only for the small share of respondents in open-ended categories. That said, re-estimating the gradient under these alternative income measures produces substantively identical results.




Figure D6: Ideological Sorting Across Real Parental Income

Notes: The figure presents mutual information M between political ideology and occupational aspirations through real parental income (CPI-adjusted to 2010 dollars), pooled across 1976–2010. Income is binned into 11 \$20,000 ranges, with the highest bin top-coded at \$200,000+. The estimates include bootstrap confidence intervals from 1,000 draws, as well as a linear fit to describe the overall income gradient. Students with “undecided” careers are excluded. Panels A and B: all three ideology categories. Panels C and D: liberals and conservatives only.

D.4. Reproducing the Income Gradient Across the Study Window

The analyses in the article pool all survey years, so the income gradient could, in principle, reflect the changing composition of survey waves rather than a stable cross-sectional pattern. To address this concern, I re-estimate the gradient separately within each five-year period of the study window, computing M between political ideology and occupational aspirations across within-year income quartiles.⁶ Figure D7 presents the results, which demonstrate that the gradient is present in every period: in each five-year window, ideological sorting rises across income quartiles, indicating that the pooled gradient is not an artifact of combining survey waves with different levels of sorting.




Figure D7: Income Gradient in Ideological Sorting Across Five-Year Periods

Notes: The figure presents mutual information M between political ideology and occupational aspirations across within-year parental income quartiles, computed separately in each period. Solid lines represent unweighted estimates, while dashed lines represent weighted estimates. Participants who are “undecided” in their career aspirations are excluded. The analyses include all ideology categories (liberal, moderate, and conservative), though substantively similar results hold when restricting the sample to liberals and conservatives only.

⁶I use income quartiles rather than deciles for these analyses to maximize power, though substantively similar results hold with deciles as well.

D.5. Students with Undecided Career Aspirations

I exclude students who report “undecided” career aspirations across all analyses in the main article. That said, the propensity to report an undecided career is not random: as Figure D8 shows, liberal students are the most likely to report undecided aspirations and conservative students the least, and undecided rates rise with parental income. Since undecidedness correlates with both ideology and income, its exclusion could, in principle, affect the income gradient. To assess this possibility, I estimate the gradient with undecided students retained as an occupational category of their own. I find that the results are substantively identical to the main results, as reported in Figure D9.




Figure D8: Share of Undecided Career Aspirations by Ideology and Income

Notes: The figure presents the share of students reporting “undecided” career aspirations, by political ideology (Panel A) and by political ideology across within-year parental income deciles (Panel B). Note that the analyses pool across 1976–2010, and as shown in Figure B1, there are time trends in the propensity to report “undecided” responses.




Figure D9: Income Gradient in Ideological Sorting, with Undecided Aspirations

Notes: The figure presents mutual information M between political ideology and occupational aspirations across household income, which I calculate as within-year income rank across 1976-2010. Students who report “undecided” career aspirations are retained as an occupational category of their own. The estimates include bootstrap confidence intervals from 1,000 draws, as well as a linear fit to describe the income gradient. Panels A and B: all three ideology categories. Panels C and D: liberals and conservatives only.

D.6. Robustness to the Composition of the Occupation Set

Figure 4 in the main article demonstrates that a small set of occupational categories carries a large share of the aggregate sorting, which raises the question of whether the income gradient depends on the presence of particular occupations. To assess this possibility, I estimate the income gradient on random subsets of occupations: in each of 1,000 runs, I draw 30 of the 40 available occupational categories, recompute M between political ideology and occupational aspirations within each income decile, and plot the final gradient. Figure D10 presents all 1,000 subset trajectories alongside the observed gradient. The rising income gradient reappears in essentially every subset.




Figure D10: Income Gradient Across Random Occupational Subsets

Notes: The figure presents mutual information M between political ideology and occupational aspirations across within-year parental income deciles, re-estimated on 1,000 random subsets of 30 occupational categories (gray lines), alongside the estimate from the full set of categories (black line). Participants who are “undecided” in their career aspirations are excluded. In each run, I randomly sample 30 occupations out of 40 categories without replacement, compute the mutual information M , and plot the income gradient within this subset.

D.7. Income Gradient Across Parental Education

Since parental income and parental education are strongly correlated, one might ask whether the income gradient documented in the main manuscript reflects material security as such, or whether it instead proxies for the educational and cultural resources of more-educated households. To test this, I stratify students into those whose parents hold a college degree and those whose parents do not, and calculate the income gradient within each group. Figure D11 presents the results from this analysis. Within both groups, I find that mutual information rises across income; at the same time, the figure shows a second pattern: at every point in the income distribution, ideological sorting is substantially higher among students from college-educated families than among their peers from non-college families. Whether this reflects the transmission of expressive or political dispositions in more-educated households, or simply that education compounds the effect of material standing, is not something this data can settle. Regardless of the reading, however, we see that the gradient survives educational stratification, with a similar income slope across both groups.




Figure D11: Income Gradient in Ideological Sorting by Parental College Attainment

Notes: The figure presents mutual information M between political ideology and occupational aspirations across household income, which I calculate as within-year income rank across 1976–2010, estimated separately within parent education levels. The parental education status is coded based on the highest degree attained by the father or the mother. The top row presents students from families without a college degree; the bottom row presents students from families with a college degree. Lines are linear fits describing the income gradient. All three ideology categories (liberal, moderate, and conservative) are included, and participants who are “undecided” in their career aspirations are excluded.

D.8. Robustness to Occupational Reproduction Across Generations

Occupational reproduction is far from evenly distributed across the income distribution. Overall, roughly 12% of students aspire to their father's or mother's occupation, but this share rises strongly with parental income: from about 6% in the bottom income quartile to about 21% in the top. It can thus be that the gradient is an artifact of this reproduction. Figure D12 finds that this is not the case, and sorting still rises across income among students who do not aspire to a parental occupation.




Figure D12: Income Gradient Excluding Occupational Reproduction

Notes: The figure presents mutual information M between political ideology and occupational aspirations across household income, which I calculate as within-year income rank across 1976–2010, comparing all students against those who do not aspire to their father's or mother's occupation. The top row includes all students; the bottom row excludes students whose intended career matches either parent's occupation. Lines are linear fits describing the income gradient. Participants who are "undecided" in their career aspirations are excluded.

D.9. Analyses Using Political Beliefs and Preferences

The income gradient I report in the article might be an artifact of socioeconomic differences in the use or comprehension of ideological labels (i.e., liberal or conservative). I reproduced this analysis using twelve issue positions as described before. Figure D13 shows that, once again, the association between each political position and occupational aspirations increases with one's parental income: all twelve estimated gradients are positive, and nearly all are distinguishable from zero. Ideology placement also produces the largest effect, consistent with its role as a summary orientation.




Figure D13: Standardized Income Gradients Across Political Issue Positions

Notes: The figure presents standardized income gradients for political self-placement and each of the twelve political issue positions. For each item, I compute the mutual information M between the item's response categories and occupational aspirations within each within-year parental income decile, normalize these values by the entropy of the item within each decile to calculate normalized mutual information, and estimate the gradient as the slope of a linear fit across the ten deciles. Standard errors for the decile-specific estimates come from 1,000 bootstrap draws and are propagated to the slopes. Horizontal bars represent 95% confidence intervals for these slope coefficients.

References

- Baldassarri, Delia, and Andrew Gelman. 2008. "Partisans Without Constraint: Political Polarization and Trends in American Public Opinion." *American Journal of Sociology* 114(2):408–46. doi:[10.1086/590649](https://doi.org/10.1086/590649).
- Eagan, Kevin, Ellen Bara Stolzenberg, Joseph J. Ramirez, Melissa C. Aragon, Maria Ramirez Suchard, and Cecilia Rios-Aguilar. 2016. *The American Freshman: Fifty Year Trends*. Los Angeles: Cooperative Institutional Research Program; the Higher Education Research Institute.
- Enke, Benjamin, Mattias K. Polborn, and Alex A. Wu. 2025. "Values as Luxury Goods and Political Behavior." *Journal of the European Economic Association* 23(5):1635–68. doi:[10.1093/jeea/jvaf014](https://doi.org/10.1093/jeea/jvaf014).
- Hinich, Melvin J., and Michael C. Munger. 1998. *Ideology and the Theory of Political Choice*. Michigan: University of Michigan Press.
- Inglehart, Ronald. 1997. *Modernization and Postmodernization: Cultural, Economic, and Political Change in 43 Societies*. Princeton: Princeton University Press.
- Kagan, Max, Justin Frake, and Reuben Hurst. 2025. "'VRscores: A New Measure and Dataset of Workforce Politics Using Voter Registrations'." SSRN, Working Paper.
- Levendusky, Matthew. 2009. *The Partisan Sort: How Liberals Became Democrats and Conservatives Became Republicans*. Chicago: University of Chicago Press.
- Mas, Alexandre, and Amanda Pallais. 2017. "Valuing Alternative Work Arrangements." *American Economic Review* 107(12):3722–59. doi:[10.1257/aer.20161500](https://doi.org/10.1257/aer.20161500).
- Rosen, Sherwin. 1986. "The Theory of Equalizing Differences." Pp. 641–92 in *Handbook of Labor Economics*. Vol. 1, edited by O. C. Ashenfelter and R. Layard. Amsterdam: North-Holland.
- Sobchuk, Oleg, and Bret Beheim. 2025. "Does Literature Evolve One Funeral at a Time?" *Proceedings of the Royal Society B: Biological Sciences* 292(2040):20242033. doi:[10.1098/rspb.2024.2033](https://doi.org/10.1098/rspb.2024.2033).